

Data Science Using Python, SAS, & R:

A Rosetta Stone for Analytical Languages

Table of Contents

Python Tutorial	5
1 Reading in Data and Basic Statistical Functions.....	5
1.1 Read in the data.....	5
a) Read the data in as a .csv file.....	6
b) Read the data in as a .xls file.	6
c) Read the data in as a .json file.	6
1.2 Find the dimensions of the data set.	6
1.3 Find basic information about the data set.....	6
1.4 Look at the first 5 (last 5) observations.	6
1.5 Calculate means of numeric variables.	7
1.6 Compute summary statistics of the data set.	7
1.7 Descriptive statistics functions applied to variables of the data set.	7
1.8 Produce a one-way table to describe the frequency of a variable.....	8
a) Produce a one-way table of a discrete variable.....	8
b) Produce a one-way table of a categorical variable.....	8
1.9 Produce a two-way table to describe the frequency of two categorical or discrete variables.	8
1.10 Select a subset of the data that meets a certain criterion.....	9
1.11 Determine the correlation between two continuous variables.	9
2 Basic Graphing and Plotting Functions.....	9
2.1 Visualize a single continuous variable by producing a histogram.	9
2.2 Visualize a single continuous variable by producing a boxplot.....	10
2.3 Visualize two continuous variables by producing a scatterplot.....	11
2.4 Visualize a relationship between two continuous variables by producing a scatterplot and a plotted line of best fit.	12
2.5 Visualize a categorical variable by producing a bar chart.	12
2.6 Visualize a continuous variable, grouped by a categorical variable, by producing side-by-side boxplots.	13
a) Simple side-by-side boxplot without color.	13

b) More advanced side-by-side boxplot with color.	14
3 Basic Data Wrangling and Manipulation	15
3.1 Create a new variable in a data set as a function of existing variables in the data set.	15
3.2 Create a new variable in a data set using if/else logic of existing variables in the data set.	15
3.3 Create new variables in a data set using mathematical functions applied to existing variables in the data set.	16
3.4 Drop variables from a data set.	16
3.5 Sort a data set by a variable.	17
a) Sort data set by a continuous variable.	17
b) Sort data set by a categorical variable.	17
3.6 Compute descriptive statistics of continuous variables, grouped by a categorical variable.	17
3.7 Add a new row to the bottom of a data set.	17
3.8 Create a user-defined function and apply it to a variable in the data set to create a new variable in the data set.	18
4 More Advanced Data Wrangling	19
4.1 Drop observations with missing information.	19
4.2 Merge two data sets together on a common variable.	19
a) First, select specific columns of a data set to create two smaller data sets.	19
b) Second, we want to merge the two smaller data sets on the common variable.	20
c) Finally, we want to check to see if the merged data set is the same as the original data set.	20
4.3 Merge two data sets together by index number only.	20
a) First, select specific columns of a data set to create two smaller data sets.	20
b) Second, we want to join the two smaller data sets.	21
c) Finally, we want to check to see if the joined data set is the same as the original data set.	21
4.4 Create a pivot table to summarize information about a data set.	21
4.5 Return all unique values from a text variable.	22
5 Preparation & Basic Regression	22
5.1 Pre-process a data set using principal component analysis.	22
5.2 Split data into training and testing data and export as a .csv file.	23
5.3 Fit a logistic regression model.	23
5.4 Fit a linear regression model.	24

6 Supervised Machine Learning	24
6.1 Fit a logistic regression model on training data and assess against testing data.	24
a) Fit a logistic regression model on training data.	24
b) Assess the model against the testing data.	25
6.2 Fit a linear regression model on training data and assess against testing data.	26
a) Fit a linear regression model on training data.	26
b) Assess the model against the testing data.	26
6.3 Fit a decision tree model on training data and assess against testing data.	26
a) Fit a decision tree classification model.	26
b) Fit a decision tree regression model.	27
6.4 Fit a random forest model on training data and assess against testing data.	28
a) Fit a random forest classification model.	28
b) Fit a random forest regression model.	29
6.5 Fit a gradient boosting model on training data and assess against testing data.	30
a) Fit a gradient boosting classification model.	30
b) Fit a gradient boosting regression model.	31
6.6 Fit an extreme gradient boosting model on training data and assess against testing data.	32
a) Fit an extreme gradient boosting classification model on training data and assess against testing data.	32
b) Fit an extreme gradient boosting regression model on training data and assess against testing data.	32
6.7 Fit a support vector model on training data and assess against testing data.	33
a) Fit a support vector classification model.	33
b) Fit a support vector regression model.	34
6.8 Fit a neural network model on training data and assess against testing data.	35
a) Fit a neural network classification model.	35
b) Fit a neural network regression model.	35
7 Unsupervised Machine Learning	36
7.1 KMeans Clustering	36
7.2 Spectral Clustering	37
7.3 Ward Hierarchical Clustering	37
7.4 DBSCAN	37
7.5 Self-organizing map	38
8 Forecasting	38

8.1 Fit an ARIMA model to a timeseries.	38
a) Plot the timeseries.....	38
b) Fit an ARIMA model and predict 2 years (24 months).....	39
8.2 Fit a Simple Exponential Smoothing model to a timeseries.....	40
a) Plot the timeseries.....	40
b) Fit a Simple Exponential Smoothing model, predict 2 years (24 months) out and plot predictions.....	41
8.3 Fit a Holt-Winters model to a timeseries.....	42
a) Plot the timeseries.....	42
b) Fit a Holt-Winters additive model, predict 2 years (24 months) out and plot predictions.....	43
9 Model Evaluation & Selection	45
9.1 Evaluate the accuracy of regression models.	45
a) Evaluation on training data.	45
b) Evaluation on testing data.	45
9.2 Evaluate the accuracy of classification models.	46
a) Evaluation on training data.	46
b) Evaluation on testing data.	46
9.3 Evaluation with cross validation.....	46
a) KFold.....	46
b) ShuffleSplit.....	47
10 Text Analytics	47
11 Deep Learning	47
Appendix.....	47
1 Built-in Python Data Types	47
2 Python Plotting Packages	48
3 Python packages used in this tutorial.....	48
Alphabetical Index	49
Array	49
Bytes & Byte arrays.....	49
Data Frame	49
Dictionary	50
List	50
Series.....	50

Sets & Frozen Sets	51
str	51

Python Tutorial

Welcome to the Python tutorial version of "Data Science Using Python, SAS, & R: A Rosetta Stone for Analytical Languages". This tutorial includes examples of common data science tasks, organized in the same way across 3 data science languages. Before beginning this tutorial, please check to make sure you have Python 3.5.2 installed (this is not required, but this is the release used to generate the following examples). Also, the following Python packages are used throughout this tutorial. You may not need all of the following packages to fit your specific needs, but they are listed below, and also in Appendix Section 2 with more detail:

pandas(0.20.2) | NumPy(1.12.1) | Matplotlib.PyPlot | seaborn(0.7.1) | re(2.2.1) | decimal(1.70) | sklearn(0.18.2) | statsmodels.api | xgboost(0.6) | pyclustering | PyFlux(0.4.15)

To install Python packages you will often need to run the following code from a terminal on your computer, and then later in a Python environment you will import the package, which is demonstrated in this tutorial:

```
pip install package_name
# or #
conda install package_name
```

Note: In Python, comments are indicated in code with a "#" character, and arrays and matrices are zero-indexed.

Now let's get started! First, you need to import several very important Python packages for data manipulation and scientific computing. The [pandas](#) package is useful for data manipulation and the [NumPy](#) package is useful for scientific computing.

```
import pandas as pd
import numpy as np
```

1 Reading in Data and Basic Statistical Functions

1.1 Read in the data.

The following demonstrate importing data into Python given 3 different file formats. The pandas package is able to read all 3 formats, as well as many others, using [Python IO tools](#).

a) Read the data in as a .csv file.

```
student = pd.read_csv('/Users/class.csv')
```

b) Read the data in as a .xls file.

Notice you must specify the file location, as well as the name of the sheet
of the .xls file you want to import

```
student_xls = pd.read_excel(open('/Users/class.xls', 'rb'),  
                             sheetname='class')
```

c) Read the data in as a .json file.

```
student_json = pd.read_json('/Users/class.json')
```

1.2 Find the dimensions of the data set.

The dimensions of a [DataFrame](#) in Python are known as an attribute of the object. Therefore, you can state the data name followed by [.shape](#) to return the dimensions of the data, with the first integer indicating the number of rows and the second indicating the number of columns.

```
print(student.shape)
```

```
## (19, 5)
```

1.3 Find basic information about the data set.

Information about a [DataFrame](#) is available by calling the [info\(\)](#) function on the data.

```
print(student.info())
```

```
## <class 'pandas.core.frame.DataFrame'>  
## RangeIndex: 19 entries, 0 to 18  
## Data columns (total 5 columns):  
## Name      19 non-null object  
## Sex       19 non-null object  
## Age       19 non-null int64  
## Height    19 non-null float64  
## Weight    19 non-null float64  
## dtypes: float64(2), int64(1), object(2)  
## memory usage: 840.0+ bytes  
## None
```

1.4 Look at the first 5 (last 5) observations.

The first 5 observations of a [DataFrame](#) are available by calling the [head\(\)](#) function on the data. By default, [head\(\)](#) returns 5 observations. To return the first *n* observations, pass the integer *n* into the function. The [tail\(\)](#) function is analogous and returns the last observations.

```
print(student.head())
```

```
##      Name Sex Age Height Weight
## 0   Alfred  M  14   69.0   112.5
## 1   Alice  F  13   56.5    84.0
## 2 Barbara  F  13   65.3    98.0
## 3   Carol  F  14   62.8   102.5
## 4   Henry  M  14   63.5   102.5
```

1.5 Calculate means of numeric variables.

The means of numeric variables of a `DataFrame` are available by calling the `mean()` function on the data.

```
print(student.mean())

## Age      13.315789
## Height   62.336842
## Weight   100.026316
## dtype: float64
```

1.6 Compute summary statistics of the data set.

Summary statistics of a `DataFrame` are available by calling the `describe()` function on the data.

```
print(student.describe())

##      Age      Height      Weight
## count  19.000000  19.000000  19.000000
## mean   13.315789  62.336842  100.026316
## std     1.492672   5.127075   22.773933
## min    11.000000  51.300000   50.500000
## 25%    12.000000  58.250000   84.250000
## 50%    13.000000  62.800000   99.500000
## 75%    14.500000  65.900000  112.250000
## max    16.000000  72.000000  150.000000
```

1.7 Descriptive statistics functions applied to variables of the data set.

Notice the subsetting of student with `[]` and the name of the variable in quotes `""`

```
print(student["Weight"].std())

## 22.773933493879046

print(student["Weight"].sum())

## 1900.5

print(student["Weight"].count())

## 19

print(student["Weight"].max())
```

```
## 150.0

print(student["Weight"].min())

## 50.5

print(student["Weight"].median())

## 99.5
```

1.8 Produce a one-way table to describe the frequency of a variable.

a) Produce a one-way table of a discrete variable.

```
# columns = "count" indicates to make the descriptive portion of the table
# the counts of each level of the index variable
print(pd.crosstab(index=student["Age"], columns="count"))
```

```
## col_0  count
## Age
## 11        2
## 12        5
## 13        3
## 14        4
## 15        4
## 16        1
```

b) Produce a one-way table of a categorical variable.

```
print(pd.crosstab(index=student["Sex"], columns="count"))
```

```
## col_0  count
## Sex
## F        9
## M       10
```

```
pd.crosstab()
```

1.9 Produce a two-way table to describe the frequency of two categorical or discrete variables.

```
# Notice the specification of a variable for the columns argument, instead
# of "count"
```

```
print(pd.crosstab(index=student["Age"], columns=student["Sex"]))
```

```
## Sex  F  M
## Age
## 11   1  1
## 12   2  3
## 13   2  1
## 14   2  2
## 15   2  2
## 16   0  1
```


`pd.crosstab()`

1.10 Select a subset of the data that meets a certain criterion.

```
females = student.query('Sex == "F"')
print(females.head())
```

```
##      Name Sex  Age  Height  Weight
## 1   Alice  F   13   56.5    84.0
## 2 Barbara  F   13   65.3    98.0
## 3   Carol  F   14   62.8   102.5
## 6    Jane  F   12   59.8    84.5
## 7   Janet  F   15   62.5   112.5
```

`query()`

1.11 Determine the correlation between two continuous variables.

```
# axis = 1 option indicates to concatenate column-wise
height_weight = pd.concat([student["Height"], student["Weight"]], axis = 1)
print(height_weight.corr(method = "pearson"))
```

```
##      Height  Weight
## Height  1.000000  0.877785
## Weight  0.877785  1.000000
```

`pd.concat() | corr()`


2 Basic Graphing and Plotting Functions

The Matplotlib [PyPlot](#) package is a standard Python package to use for plotting. For more information on other Python plotting packages, please see the Appendix Section 2.

```
import matplotlib.pyplot as plt
```

2.1 Visualize a single continuous variable by producing a histogram.

```
# Notice the labeling of the axes
plt.hist(student["Weight"], bins=[40,60,80,100,120,140,160])
plt.xlabel('Weight')
plt.ylabel('Frequency')
plt.show()
```




Output:

2.2 Visualize a single continuous variable by producing a boxplot.

```
# showmeans=True tells Python to plot the mean of the variable on the boxplot  
plt.boxplot(student["Weight"], showmeans=True)
```

```
# prevents Python from printing a "1" at the bottom of the boxplot  
plt.xticks([])
```


```
plt.ylabel('Weight')  
plt.show()
```



Output:

2.3 Visualize two continuous variables by producing a scatterplot.

```
# Notice here you specify the x variable, followed by the y variable
plt.scatter(student["Height"], student["Weight"])
plt.xlabel("Height")
plt.ylabel("Weight")
plt.show()
```




Output:

2.4 Visualize a relationship between two continuous variables by producing a scatterplot and a plotted line of best fit.

```
x = student["Height"]
y = student["Weight"]

# np.polyfit() models Weight as a function of Height and returns the
# parameters
m, b = np.polyfit(x, y, 1)
plt.scatter(x, y)

# plt.text() prints the equation of the line of best fit, with the first two
# arguments specifying the x and y locations of the text, respectively
# "%f" indicates to print a floating point number, that is specified
following
# the string and a "%" character
plt.text(51, 140, "Line: y = %f x + %f"% (m,b))
plt.plot(x, m*x + b)
plt.xlabel("Height")
plt.ylabel("Weight")
plt.show()
```



Output:

`np.polyfit()`


2.5 Visualize a categorical variable by producing a bar chart.

```
# Get the counts of Sex
counts = pd.crosstab(index=student["Sex"], columns="count")
```

```
# len() returns the number of categories of Sex (2)
# np.arange() creates a vector of the specified length
num = np.arange(len(counts))

# alpha = 0.5 changes the transparency of the bars
plt.bar(num, counts["count"], align='center', alpha=0.5)

# Set the xticks to be the indices of counts
plt.xticks(num, counts.index)
plt.xlabel("Sex")
plt.ylabel("Frequency")
plt.show()
```



Output:

`np.arange()`


2.6 Visualize a continuous variable, grouped by a categorical variable, by producing side-by-side boxplots.

a) Simple side-by-side boxplot without color.

```
# Subset data set to return only female weights, and then only male weights
Weight_F = np.array(student.query('Sex == "F"')['Weight'])
Weight_M = np.array(student.query('Sex == "M"')['Weight'])
Weights = [Weight_F, Weight_M]
```

```
# PyPlot automatically plots the two weights side-by-side since Weights
```

```
# is a 2D array
plt.boxplot(Weights, showmeans=True, labels=('F', 'M'))
plt.xlabel('Sex')
plt.ylabel('Weight')
plt.show()
```

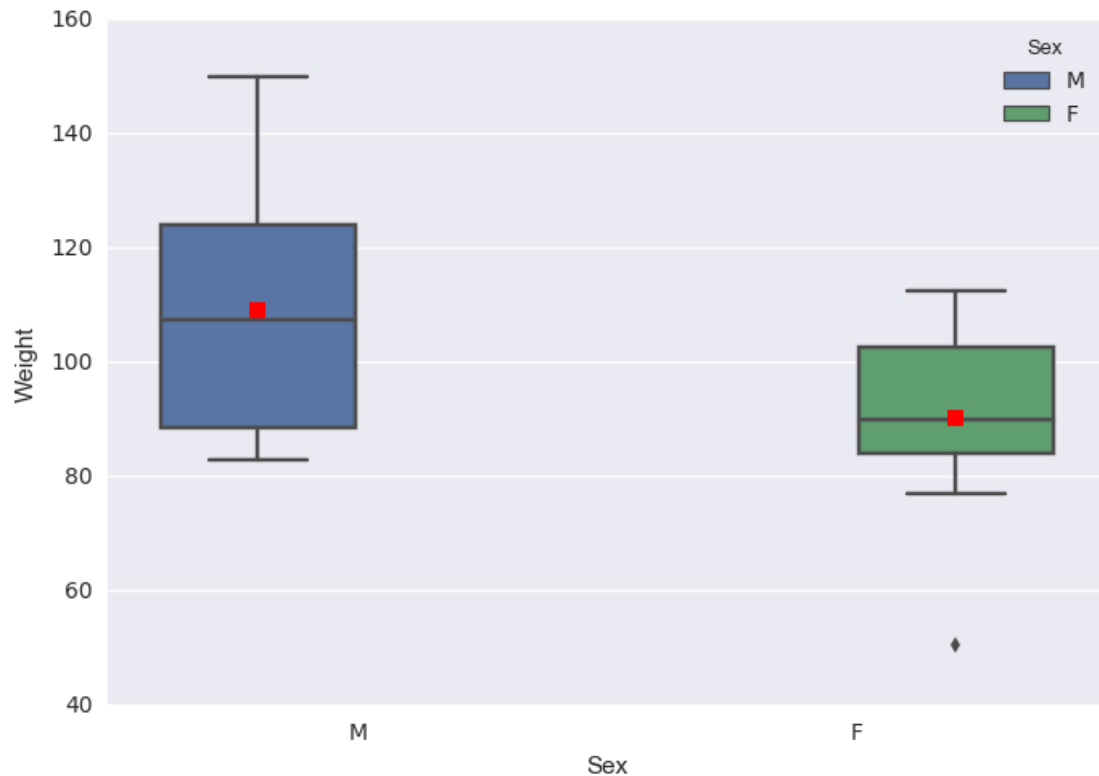


Output:

```
np.array()
```

b) More advanced side-by-side boxplot with color.

```
import seaborn as sns
sns.boxplot(x="Sex", y="Weight", hue="Sex", data = student, showmeans=True)
sns.plt.show()
```



Output:
seaborn

3 Basic Data Wrangling and Manipulation

3.1 Create a new variable in a data set as a function of existing variables in the data set.

```
# Notice here how you can create the BMI column in the data set
# just by naming it
student["BMI"] = student["Weight"] / student["Height"]**2 * 703
print(student.head())
```

##	Name	Sex	Age	Height	Weight	BMI
## 0	Alfred	M	14	69.0	112.5	16.611531
## 1	Alice	F	13	56.5	84.0	18.498551
## 2	Barbara	F	13	65.3	98.0	16.156788
## 3	Carol	F	14	62.8	102.5	18.270898
## 4	Henry	M	14	63.5	102.5	17.870296

3.2 Create a new variable in a data set using if/else logic of existing variables in the data set.

```
# Notice the use of the np.where() function for a single condition
student["BMI Class"] = np.where(student["BMI"] < 19.0, "Underweight",
```

```

                                "Healthy")
print(student.head())

```

	Name	Sex	Age	Height	Weight	BMI	BMI Class
## 0	Alfred	M	14	69.0	112.5	16.611531	Underweight
## 1	Alice	F	13	56.5	84.0	18.498551	Underweight
## 2	Barbara	F	13	65.3	98.0	16.156788	Underweight
## 3	Carol	F	14	62.8	102.5	18.270898	Underweight
## 4	Henry	M	14	63.5	102.5	17.870296	Underweight

```
np.where()
```

3.3 Create new variables in a data set using mathematical functions applied to existing variables in the data set.

Using the `np.log()`, `np.exp()`, `np.sqrt()`, `np.where()`, and `np.abs()` functions.

```

student["LogWeight"] = np.log(student["Weight"])
student["ExpAge"] = np.exp(student["Age"])
student["SqrtHeight"] = np.sqrt(student["Height"])
student["BMI Neg"] = np.where(student["BMI"] < 19.0, -student["BMI"],
                             student["BMI"])
student["BMI Pos"] = np.abs(student["BMI Neg"])

# Create a Boolean variable
student["BMI Check"] = (student["BMI Pos"] == student["BMI"])

```

	Name	Sex	Age	Height	Weight	BMI	BMI Class	LogWeight	\
## 0	Alfred	M	14	69.0	112.5	16.611531	Underweight	4.722953	
## 1	Alice	F	13	56.5	84.0	18.498551	Underweight	4.430817	
## 2	Barbara	F	13	65.3	98.0	16.156788	Underweight	4.584967	
## 3	Carol	F	14	62.8	102.5	18.270898	Underweight	4.629863	
## 4	Henry	M	14	63.5	102.5	17.870296	Underweight	4.629863	

	ExpAge	SqrtHeight	BMI Neg	BMI Pos	BMI Check
## 0	1.202604e+06	8.306624	-16.611531	16.611531	True
## 1	4.424134e+05	7.516648	-18.498551	18.498551	True
## 2	4.424134e+05	8.080842	-16.156788	16.156788	True
## 3	1.202604e+06	7.924645	-18.270898	18.270898	True
## 4	1.202604e+06	7.968689	-17.870296	17.870296	True

3.4 Drop variables from a data set.

```

# axis = 1 indicates to drop columns instead of rows
student = student.drop(["LogWeight", "ExpAge", "SqrtHeight", "BMI Neg",
                       "BMI Pos", "BMI Check"], axis = 1)
print(student.head())

```

	Name	Sex	Age	Height	Weight	BMI	BMI Class
## 0	Alfred	M	14	69.0	112.5	16.611531	Underweight
## 1	Alice	F	13	56.5	84.0	18.498551	Underweight


```
## 2  Barbara  F   13   65.3   98.0  16.156788  Underweight
## 3    Carol  F   14   62.8  102.5  18.270898  Underweight
## 4    Henry  M   14   63.5  102.5  17.870296  Underweight
```

drop()

3.5 Sort a data set by a variable.

a) Sort data set by a continuous variable.

```
# Notice kind="mergesort" which indicates to use a stable sorting
# algorithm
```

```
student = student.sort_values(by="Age", kind="mergesort")
print(student.head())
```

```
##      Name Sex Age Height Weight      BMI BMI Class
## 10  Joyce  F  11   51.3   50.5  13.490001 Underweight
## 17 Thomas  M  11   57.5   85.0  18.073346 Underweight
## 5   James  M  12   57.3   83.0  17.771504 Underweight
## 6   Jane   F  12   59.8   84.5  16.611531 Underweight
## 9   John   M  12   59.0   99.5  20.094369 Healthy
```

b) Sort data set by a categorical variable.

```
student = student.sort_values(by="Sex", kind="mergesort")
# Notice that the data is now sorted first by Sex and then within Sex by Age
print(student.head())
```

```
##      Name Sex Age Height Weight      BMI BMI Class
## 10  Joyce  F  11   51.3   50.5  13.490001 Underweight
## 6   Jane   F  12   59.8   84.5  16.611531 Underweight
## 12 Louise  F  12   56.3   77.0  17.077695 Underweight
## 1   Alice  F  13   56.5   84.0  18.498551 Underweight
## 2  Barbara  F  13   65.3   98.0  16.156788 Underweight
```

sort_values()

3.6 Compute descriptive statistics of continuous variables, grouped by a categorical variable.

```
print(student.groupby(by="Sex").mean())
```

```
##      Age      Height      Weight      BMI
## Sex
## F    13.222222  60.588889   90.111111  17.051039
## M    13.400000  63.910000  108.950000  18.594243
```

groupby()

3.7 Add a new row to the bottom of a data set.

```
# Look at the tail of the data currently
print(student.tail())
```

```
##      Name Sex Age Height Weight      BMI BMI Class
## 0    Alfred M  14   69.0  112.5  16.611531 Underweight
## 4     Henry M  14   63.5  102.5  17.870296 Underweight
## 16   Ronald M  15   67.0  133.0  20.828470   Healthy
## 18  William M  15   66.5  112.0  17.804511 Underweight
## 14   Philip M  16   72.0  150.0  20.341435   Healthy

student = student.append({'Name': 'Jane', 'Sex': 'F', 'Age': 14, 'Height': 56.3,
                          'Weight': 77.0, 'BMI': 17.077695,
                          'BMI Class': 'Underweight'},
                          ignore_index=True)
```

```
# Notice the change in the indices because of the ignore_index=True option
# which allows for a Series, or one-dimensional DataFrame, to be appended
# to an existing DataFrame
```

```
##      Name Sex Age Height Weight      BMI BMI Class
## 15   Henry M  14   63.5  102.5  17.870296 Underweight
## 16   Ronald M  15   67.0  133.0  20.828470   Healthy
## 17  William M  15   66.5  112.0  17.804511 Underweight
## 18   Philip M  16   72.0  150.0  20.341435   Healthy
## 19    Jane F  14   56.3   77.0  17.077695 Underweight
```

`append()`

3.8 Create a user-defined function and apply it to a variable in the data set to create a new variable in the data set.

```
def toKG(lb):
    return (0.45359237 * lb)

student["Weight KG"] = student["Weight"].apply(toKG)
print(student.head())
```

```
##      Name Sex Age Height Weight      BMI BMI Class Weight KG
## 0    Joyce F  11   51.3   50.5  13.490001 Underweight  22.906415
## 1     Jane F  12   59.8   84.5  16.611531 Underweight  38.328555
## 2   Louise F  12   56.3   77.0  17.077695 Underweight  34.926612
## 3    Alice F  13   56.5   84.0  18.498551 Underweight  38.101759
## 4  Barbara F  13   65.3   98.0  16.156788 Underweight  44.452052
```

`apply()` | user-defined functions

4 More Advanced Data Wrangling

4.1 Drop observations with missing information.

```
# Notice the use of the fish data set because it has some missing
# observations
fish = pd.read_csv('/Users/fish.csv')

# First sort by Weight, requesting those with NA for Weight first
fish = fish.sort_values(by='Weight', kind='mergesort', na_position='first')
print(fish.head())
```

	Species	Weight	Length1	Length2	Length3	Height	Width
## 13	Bream	NaN	29.5	32.0	37.3	13.9129	5.0728
## 40	Roach	0.0	19.0	20.5	22.8	6.4752	3.3516
## 72	Perch	5.9	7.5	8.4	8.8	2.1120	1.4080
## 145	Smelt	6.7	9.3	9.8	10.8	1.7388	1.0476
## 147	Smelt	7.0	10.1	10.6	11.6	1.7284	1.1484

--

```
new_fish = fish.dropna()
print(new_fish.head())
```

	Species	Weight	Length1	Length2	Length3	Height	Width
## 40	Roach	0.0	19.0	20.5	22.8	6.4752	3.3516
## 72	Perch	5.9	7.5	8.4	8.8	2.1120	1.4080
## 145	Smelt	6.7	9.3	9.8	10.8	1.7388	1.0476
## 147	Smelt	7.0	10.1	10.6	11.6	1.7284	1.1484
## 146	Smelt	7.5	10.0	10.5	11.6	1.9720	1.1600

`dropna()`

4.2 Merge two data sets together on a common variable.

a) First, select specific columns of a data set to create two smaller data sets.

```
# Notice the use of the student data set again, however we want to reload it
# without the changes we've made previously
student = pd.read_csv('/Users/class.csv')
student1 = pd.concat([student["Name"], student["Sex"], student["Age"]],
                     axis = 1)
print(student1.head())
```

	Name	Sex	Age
## 0	Alfred	M	14
## 1	Alice	F	13
## 2	Barbara	F	13
## 3	Carol	F	14
## 4	Henry	M	14

--

```
student2 = pd.concat([student["Name"], student["Height"], student["Weight"]],
                     axis = 1)
print(student2.head())
```

```
##      Name  Height  Weight
## 0  Alfred    69.0   112.5
## 1   Alice    56.5    84.0
## 2 Barbara    65.3    98.0
## 3   Carol    62.8   102.5
## 4   Henry    63.5   102.5
```

b) Second, we want to merge the two smaller data sets on the common variable.

```
new = pd.merge(student1, student2, on="Name")
print(new.head())
```

```
##      Name Sex  Age  Height  Weight
## 0  Alfred  M   14    69.0   112.5
## 1   Alice  F   13    56.5    84.0
## 2 Barbara  F   13    65.3    98.0
## 3   Carol  F   14    62.8   102.5
## 4   Henry  M   14    63.5   102.5
```

```
pd.merge()
```

c) Finally, we want to check to see if the merged data set is the same as the original data set.

```
print(student.equals(new))
```

```
## True
```

```
equals()
```

4.3 Merge two data sets together by index number only.

a) First, select specific columns of a data set to create two smaller data sets.

```
newstudent1 = pd.concat([student["Name"], student["Sex"], student["Age"]],
                       axis = 1)
print(newstudent1.head())
```

```
##      Name Sex  Age
## 0  Alfred  M   14
## 1   Alice  F   13
## 2 Barbara  F   13
## 3   Carol  F   14
## 4   Henry  M   14
```

```
--
```

```
newstudent2 = pd.concat([student["Height"], student["Weight"]], axis = 1)
print(newstudent2.head())
```

```
##      Height  Weight
## 0      69.0    112.5
## 1      56.5     84.0
## 2      65.3     98.0
## 3      62.8    102.5
## 4      63.5    102.5
```

b) Second, we want to join the two smaller data sets.

```
new2 = newstudent1.join(newstudent2)
print(new2.head())
```

```
##      Name Sex  Age  Height  Weight
## 0  Alfred  M   14    69.0    112.5
## 1  Alice   F   13    56.5     84.0
## 2 Barbara  F   13    65.3     98.0
## 3  Carol   F   14    62.8    102.5
## 4  Henry   M   14    63.5    102.5
```

join()

c) Finally, we want to check to see if the joined data set is the same as the original data set.

```
print(student.equals(new2))
```

```
## True
```

4.4 Create a pivot table to summarize information about a data set.

```
# Notice we are using a new data set that needs to be read into the
# environment
```

```
price = pd.read_csv('/Users/price.csv')
```

```
# The following code is used to remove the "," and "$" characters from
# the ACTUAL colum so that the values can be summed
```

```
from re import sub
from decimal import Decimal
def trim_money(money):
    return(float(Decimal(sub(r'^\d.', '', money))))
```

```
price["REVENUE"] = price["ACTUAL"].apply(trim_money)
table = pd.pivot_table(price, index=["COUNTRY", "STATE", "PRODTYPE",
                                     "PRODUCT"], values="REVENUE",
                       aggfunc=np.sum)
print(table.head())
```

```
##                                     REVENUE
## COUNTRY STATE      PRODTYPE  PRODUCT
## Canada  British Columbia FURNITURE BED      197706.6
```

```
##                                SOFA      216282.6
##                                OFFICE    CHAIR    200905.2
##                                DESK      186262.2
##                                Ontario  FURNITURE BED    194493.6
```

Note: `pd.pivot_table()` is similar to the `pd.pivot()` function.

[re](#) | [Decimal](#)

4.5 Return all unique values from a text variable.

```
print(np.unique(price["STATE"]))

## ['Baja California Norte' 'British Columbia' 'California' 'Campeche'
##  'Colorado' 'Florida' 'Illinois' 'Michoacan' 'New York' 'North Carolina'
##  'Nuevo Leon' 'Ontario' 'Quebec' 'Saskatchewan' 'Texas' 'Washington']
```

[np.unique\(\)](#)

The following sections focus on the Python [sklearn](#) package. Also, in the following sections, several data set will be used more than once for prediction and modeling. Often, they will be re-read into the environment so we are always going back to the original, raw data.

5 Preparation & Basic Regression

5.1 Pre-process a data set using principal component analysis.

```
# Notice we are using a new data set that needs to be read into the
# environment
```

```
iris = pd.read_csv('/Users/iris.csv')
features = iris.drop(["Target"], axis = 1)
```

```
from sklearn import preprocessing
features_scaled = preprocessing.scale(features.as_matrix())
```

```
from sklearn.decomposition import PCA
```

```
pca = PCA(n_components = 4)
pca = pca.fit(features_scaled)
print(np.transpose(pca.components_))
```

```
## [[ 0.52237162  0.37231836 -0.72101681 -0.26199559]
##  [-0.26335492  0.92555649  0.24203288  0.12413481]
##  [ 0.58125401  0.02109478  0.14089226  0.80115427]
##  [ 0.56561105  0.06541577  0.6338014  -0.52354627]]
```

[preprocessing](#) | [PCA](#) | [np.transpose\(\)](#)

5.2 Split data into training and testing data and export as a .csv file.

```
from sklearn.model_selection import train_test_split

target = iris["Target"]

# The following code splits the iris data set into 70% train and 30% test
X_train, X_test, Y_train, Y_test = train_test_split(features, target,
                                                    test_size = 0.3,
                                                    random_state = 29)

train_x = pd.DataFrame(X_train)
train_y = pd.DataFrame(Y_train)
test_x = pd.DataFrame(X_test)
test_y = pd.DataFrame(Y_test)

train = pd.concat([train_x, train_y], axis = 1)
test = pd.concat([test_x, test_y], axis = 1)

train.to_csv('/Users/iris_train_Python.csv', index = False)
test.to_csv('/Users/iris_test_Python.csv', index = False)
```

[train_test_split\(\)](#) | [to_csv\(\)](#)

5.3 Fit a logistic regression model.

```
# Notice we are using a new data set that needs to be read into the
# environment
tips = pd.read_csv('/Users/tips.csv')

# The following code is used to determine if the individual left more
# than a 15% tip
tips["fifteen"] = 0.15 * tips["total_bill"]
tips["greater15"] = np.where(tips["tip"] > tips["fifteen"], 1, 0)

import statsmodels.api as sm

# Notice the syntax of greater15 as a function of total_bill
logreg = sm.formula.glm("greater15 ~ total_bill",
                        family=sm.families.Binomial(),
                        data=tips).fit()

print(logreg.summary())
```

```
##                               Generalized Linear Model Regression Results                               ##
=====
## Dep. Variable:                 greater15    No. Observations:                 244
## Model:                       GLM           Df Residuals:                     242
## Model Family:                 Binomial      Df Model:                         1
## Link Function:                logit         Scale:                           1.0
## Method:                       IRLS         Log-Likelihood:                   -156.87
## Date:                         Mon, 03 Jul 2017    Deviance:                       313.74
```

```
## Time: 09:16:27 Pearson chi2: 247.
## No. Iterations: 4
##
```

##	coef	std err	z	P> z	[0.025	0.975]
##						
## Intercept	1.6477	0.355	4.646	0.000	0.953	2.343
## total_bill	-0.0725	0.017	-4.319	0.000	-0.105	-0.040
##						

A logistic regression model can be implemented using `sklearn`, however `statsmodels.api` provides a helpful summary about the model, so it is preferable for this example.

5.4 Fit a linear regression model.

```
# Fit a linear regression model of tip by total_bill
from sklearn.linear_model import LinearRegression

# If your data has one feature, you need to reshape the 1D array
linreg = LinearRegression()
linreg.fit(tips["total_bill"].values.reshape(-1,1), tips["tip"])
print(linreg.coef_)
print(linreg.intercept_)

## [ 0.10502452]
## 0.920269613555
```

LinearRegression

6 Supervised Machine Learning

6.1 Fit a logistic regression model on training data and assess against testing data.

a) Fit a logistic regression model on training data.

```
# Notice we are using new data sets that need to be read into the environment
train = pd.read_csv('/Users/tips_train.csv')
test = pd.read_csv('/Users/tips_test.csv')

train["fifteen"] = 0.15 * train["total_bill"]
train["greater15"] = np.where(train["tip"] > train["fifteen"], 1, 0)
test["fifteen"] = 0.15 * test["total_bill"]
test["greater15"] = np.where(test["tip"] > test["fifteen"], 1, 0)

logreg = sm.formula.glm("greater15 ~ total_bill",
                        family=sm.families.Binomial(),
```



```

data=train).fit()
print(logreg.summary())

```

```

##                               Generalized Linear Model Regression Results                               ##
=====
## Dep. Variable:                greater15    No. Observations:                195
## Model:                      GLM           Df Residuals:                  193
## Model Family:               Binomial      Df Model:                      1
## Link Function:              logit         Scale:                      1.0
## Method:                     IRLS         Log-Likelihood:                 -125.29
## Date:                       Mon, 03 Jul 2017    Deviance:                      250.58
## Time:                       09:16:30         Pearson chi2:                  197.
## No. Iterations:              4
=====
##               coef      std err          z      P>|z|      [0.025   0.975]
## -----
## Intercept      1.6461      0.395       4.172     0.000       0.873    2.420
## total_bill    -0.0706      0.018      -3.820     0.000      -0.107   -0.034
##
=====

```

b) Assess the model against the testing data.

```

# Prediction on testing data
predictions = logreg.predict(test["total_bill"])
predY = np.where(predictions < 0.5, 0, 1)

# If the prediction probability is less than 0.5, classify this as a 0
# and otherwise classify as a 1. This isn't the best method -- a better
# method would be randomly assigning a 0 or 1 when a probability of 0.5
# occurs, but this insures that results are consistent

# Determine how many were correctly classified
Results = np.where(predY == test["greater15"], "Correct", "Wrong")
print(pd.crosstab(index=Results, columns="count"))

## col_0    count
## row_0
## Correct    34
## Wrong      15

```

A logistic regression model can be implemented using [sklearn](#), however [statsmodels.api](#) provides a helpful summary about the model, so it is preferable for this example.

6.2 Fit a linear regression model on training data and assess against testing data.

a) Fit a linear regression model on training data.

```
# Notice we are using new data sets that need to be read into the environment
train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

# Fit a linear regression model
linreg = LinearRegression()
linreg.fit(train.drop(["Target"], axis = 1), train["Target"])
print(linreg.coef_)
print(linreg.intercept_)

## [ -8.56336900e-02  4.60343577e-02  3.64131905e-02  3.24796064e+00
## -1.48729382e+01  3.57686873e+00 -8.70316831e-03 -1.36890461e+00
##  3.13120107e-01 -1.28815611e-02 -9.76900124e-01  1.13257346e-02
## -5.26715028e-01]
## 36.1081957809
```

b) Assess the model against the testing data.

```
# Prediction on testing data
prediction = pd.DataFrame()
prediction["predY"] = linreg.predict(test.drop(["Target"], axis = 1))

# Determine mean squared error
prediction["sq_diff"] = (prediction["predY"] - test["Target"])**2
print(np.mean(prediction["sq_diff"]))

## 17.771307958891672
```

LinearRegression

6.3 Fit a decision tree model on training data and assess against testing data.

a) Fit a decision tree classification model.

i) Fit a decision tree classification model on training data and determine variable importance.

```
# Notice we are using new data sets that need to be read into the environment
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

from sklearn.tree import DecisionTreeClassifier

# random_state is used to specify a seed for a random integer so that the
# results are reproducible
treeMod = DecisionTreeClassifier(random_state=29)
```

```

treeMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Determine variable importance
var_import = treeMod.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0: 'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                     ascending = False)

print(var_import.head())

##      Importance
## 23      0.692681
## 27      0.158395
## 21      0.044384
## 11      0.029572
## 24      0.020485

```

ii) Assess the model against the testing data.

```

# Prediction on testing data
predY = treeMod.predict(test.drop(["Target"], axis = 1))

# Determine how many were correctly classified
Results = np.where(test["Target"] == predY, "Correct", "Wrong")
print(pd.crosstab(index=Results, columns="count"))

## col_0    count
## row_0
## Correct    161
## Wrong      10

```

DecisionTreeClassifier

b) Fit a decision tree regression model.

i) Fit a decision tree regression model on training data and determine variable importance.

```

train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

from sklearn.tree import DecisionTreeRegressor

treeMod = DecisionTreeRegressor(random_state=29)
treeMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Determine variable importance
var_import = treeMod.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0: 'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                     ascending = False)

print(var_import.head())

```

```
##      Importance
## 5      0.573257
## 12     0.203677
## 7      0.103939
## 4      0.041467
## 0      0.033798
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
prediction = pd.DataFrame()
prediction["predY"] = treeMod.predict(test.drop(["Target"], axis = 1))

# Determine mean squared error
prediction["sq_diff"] = (prediction["predY"] - test["Target"])**2
print(np.mean(prediction["sq_diff"]))

## 23.866842105263157
```

DecisionTreeRegressor

6.4 Fit a random forest model on training data and assess against testing data.

a) Fit a random forest classification model.

i) Fit a random forest classification model on training data and determine variable importance.

```
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

from sklearn.ensemble import RandomForestClassifier

rfMod = RandomForestClassifier(random_state=29)
rfMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Determine variable importance
var_import = rfMod.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0: 'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                     ascending = False)

print(var_import.head())

##      Importance
## 27      0.271730
## 13      0.120096
## 23      0.101971
## 20      0.076891
## 6       0.066836
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
predY = rfMod.predict(test.drop(["Target"], axis = 1))

# Determine how many were correctly classified
Results = np.where(test["Target"] == predY, "Correct", "Wrong")
print(pd.crosstab(index=Results, columns="count"))

## col_0    count
## row_0
## Correct    165
## Wrong       6
```

RandomForestClassifier

b) Fit a random forest regression model.

i) Fit a random forest regression model on training data and determine variable importance.

```
train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

from sklearn.ensemble import RandomForestRegressor

rfMod = RandomForestRegressor(random_state=29)
rfMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Determine variable importance
var_import = rfMod.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0:'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                     ascending = False)

print(var_import.head())

##      Importance
## 5      0.412012
## 12     0.392795
## 7      0.079462
## 0      0.041911
## 9      0.016374
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
prediction = pd.DataFrame()
prediction["predY"] = rfMod.predict(test.drop(["Target"], axis = 1))

# Determine mean squared error
prediction["sq_diff"] = (test["Target"] - prediction["predY"])**2
print(prediction["sq_diff"].mean())
```

```
## 13.25032631578948
```

RandomForestRegressor

6.5 Fit a gradient boosting model on training data and assess against testing data.

a) Fit a gradient boosting classification model.

i) Fit a gradient boosting classification model on training data and determine variable importance.

```
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

from sklearn.ensemble import GradientBoostingClassifier

# n_estimators = total number of trees to fit which is analogous to the
# number of iterations
# learning_rate = shrinkage or step-size reduction, where a lower
# learning rate requires more iterations
gbMod = GradientBoostingClassifier(random_state = 29, learning_rate = .01,
                                   n_estimators = 2500)
gbMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Determine variable importance
var_import = gbMod.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0: 'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                   ascending = False)

print(var_import.head())

##      Importance
## 23      0.099054
## 27      0.088744
## 7       0.062735
## 21      0.043547
## 14      0.042328
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
predY = gbMod.predict(test.drop(["Target"], axis = 1))

# Determine how many were correctly classified
Results = np.where(test["Target"] == predY, "Correct", "Wrong")
print(pd.crosstab(index=Results, columns="count"))

## col_0    count
## row_0
```

```
## Correct      164
## Wrong        7
```

GradientBoostingClassifier

b) Fit a gradient boosting regression model.

i) Fit a gradient boosting regression model on training data and determine variable importance.

```
train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

from sklearn.ensemble import GradientBoostingRegressor

gbMod = GradientBoostingRegressor(random_state = 29, learning_rate = .01,
                                   n_estimators = 2500)
gbMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Determine variable importance
var_import = gbMod.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0: 'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                     ascending = False)

print(var_import.head())

##      Importance
## 5      0.166179
## 12     0.154570
## 0      0.127526
## 11     0.124045
## 6      0.115200
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
prediction = pd.DataFrame()
prediction["predY"] = gbMod.predict(test.drop(["Target"], axis = 1))

# Determine mean squared error
prediction["sq_diff"] = (test["Target"] - prediction["predY"])**2

## 9.416022842108923
```

GradientBoostingRegressor

6.6 Fit an extreme gradient boosting model on training data and assess against testing data.

a) Fit an extreme gradient boosting classification model on training data and assess against testing data.

i) Fit an extreme gradient boosting classification model on training data.

```
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

from xgboost import XGBClassifier

# Fit XGBClassifier model on training data
xgbMod = XGBClassifier(seed = 29, learning_rate = 0.01,
                       n_estimators = 2500)
xgbMod.fit(train.drop(["Target"], axis = 1), train["Target"])
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
predY = xgbMod.predict(test.drop(["Target"], axis = 1))

# Determine how many were correctly classified
Results = np.where(test["Target"] == predY, "Correct", "Wrong")
print(pd.crosstab(index=Results, columns="count"))
```

```
## col_0    count
## row_0
## Correct    165
## Wrong       6
```

xgboost

b) Fit an extreme gradient boosting regression model on training data and assess against testing data.

i) Fit an extreme gradient boosting regression model on training data.

```
train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

from xgboost import XGBRegressor

# Fit XGBRegressor model on training data
xgbMod = XGBRegressor(seed = 29, learning_rate = 0.01,
                      n_estimators = 2500)
xgbMod.fit(train.drop(["Target"], axis = 1), train["Target"])
```


ii) Assess the model against the testing data.

```
# Prediction on testing data
prediction = pd.DataFrame()
prediction["predY"] = xgbMod.predict(test.drop(["Target"], axis = 1))

# Determine mean squared error
prediction["sq_diff"] = (test["Target"] - prediction["predY"])**2
print(prediction["sq_diff"].mean())

## 9.658108024646909
```

xgboost

6.7 Fit a support vector model on training data and assess against testing data.

a) Fit a support vector classification model.

i) Fit a support vector classification model on training data.

```
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

# First we should scale the data, since R does
from sklearn.preprocessing import StandardScaler

train_features = train.drop(["Target"], axis = 1)
scaler = StandardScaler().fit(np.array(train_features))
train_scaled = scaler.transform(np.array(train_features))
train_scaled = pd.DataFrame(train_scaled)
train_scaled["Target"] = train["Target"]

test_features = test.drop(["Target"], axis = 1)
scaler = StandardScaler().fit(np.array(test_features))
test_scaled = scaler.transform(np.array(test_features))
test_scaled = pd.DataFrame(test_scaled)
test_scaled["Target"] = test["Target"]

# Fit a support vector classification model
from sklearn.svm import SVC
svMod = SVC(random_state = 29, kernel = 'linear')
svMod.fit(train_scaled.drop(["Target"], axis = 1),
          train_scaled["Target"])
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
prediction = svMod.predict(test_scaled.drop(["Target"], axis = 1))

# Determine how many were correctly classified
```

```
Results = np.where(test_scaled["Target"] == prediction, "Correct", "Wrong")
print(pd.crosstab(index=Results, columns="count"))
```

```
## col_0    count
## row_0
## Correct   164
## Wrong      7
```

[preprocessing](#) | [SVC](#)

b) Fit a support vector regression model.

i) Fit a support vector regression model on training data.

```
train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler().fit(np.array(train))
train_scaled = scaler.transform(np.array(train))
train_scaled = pd.DataFrame(train_scaled)
```

```
scaler = StandardScaler().fit(np.array(test))
test_scaled = scaler.transform(np.array(test))
test_scaled = pd.DataFrame(test_scaled)
```

```
# Fit a support vector regression model
from sklearn.svm import SVR
svMod = SVR()
svMod.fit(train_scaled.drop([13], axis = 1), train_scaled[13])
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
prediction = pd.DataFrame()
# 13 is the name of the Target variable
prediction["predY"] = svMod.predict(test_scaled.drop([13], axis = 1))
```

```
# Determine mean squared error
prediction["sq_diff"] = (test_scaled[13] - prediction["predY"])**2
print(prediction["sq_diff"].mean())
```

```
## 0.15397129809389054
```

[preprocessing](#) | [SVR](#)

6.8 Fit a neural network model on training data and assess against testing data.

a) Fit a neural network classification model.

i) Fit a neural network classification model on training data.

```
# Notice we are using new data sets
train = pd.read_csv('/Users/digits_train.csv')
test = pd.read_csv('/Users/digits_test.csv')

# Fit a neural network classification model on training data
from sklearn.neural_network import MLPClassifier
nnMod = MLPClassifier(max_iter = 200, hidden_layer_sizes=(100,),
                      random_state = 29)
```

ii) Assess the model against the testing data.

```
# Prediction on testing data
predY = nnMod.predict(test.drop(["Target"], axis = 1))

# Determine how many were correctly classified

from sklearn.metrics import confusion_matrix

print(confusion_matrix(test["Target"], predY))

## [[57  0  0  0  1  0  0  0  0  0]
##   [ 0 57  0  0  0  0  0  0  1  0]
##   [ 0  0 58  0  0  0  0  0  0  0]
##   [ 0  0  0 58  0  1  0  0  0  0]
##   [ 0  0  0  0 52  0  1  0  1  0]
##   [ 0  0  0  0  1 56  0  1  1  0]
##   [ 0  0  0  0  0  0 41  0  0  0]
##   [ 0  0  0  0  1  0  0 49  0  1]
##   [ 0  1  0  1  0  0  0  0 43  0]
##   [ 0  1  0  0  0  1  0  0  2 53]]
```

MLPClassifier | confusion_matrix()

b) Fit a neural network regression model.

i) Fit a neural network regression model on training data.

```
train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

# Scale input data
from sklearn.preprocessing import StandardScaler

train_features = train.drop(["Target"], axis = 1)
scaler = StandardScaler().fit(np.array(train_features))
```

```

train_scaled = scaler.transform(np.array(train_features))
train_scaled = pd.DataFrame(train_scaled)
train_scaled["Target"] = train["Target"]

test_features = test.drop(["Target"], axis = 1)
scaler = StandardScaler().fit(np.array(test_features))
test_scaled = scaler.transform(np.array(test_features))
test_scaled = pd.DataFrame(test_scaled)
test_scaled["Target"] = test["Target"]

# Fit neural network regression model, dividing target by 50 for scaling
from sklearn.neural_network import MLPRegressor
nnMod = MLPRegressor(max_iter = 250, random_state = 29, solver = 'lbfgs')
nnMod = nnMod.fit(train_scaled, train["Target"] / 50)

ii) Assess the model against testing data.
# Prediction on testing data, remembering to multiply by 50
prediction = pd.DataFrame()
prediction["predY"] = nnMod.predict(test_scaled)*50

# Determine mean squared error
prediction["sq_diff"] = (test["Target"] - prediction["predY"])**2
print(prediction["sq_diff"].mean())

## 1.420655650498798

```

[preprocessing](#) | [MLPRegressor](#)

7 Unsupervised Machine Learning

7.1 KMeans Clustering

```

iris = pd.read_csv('/Users/iris.csv')
iris["Species"] = np.where(iris["Target"] == 0, "Setosa",
                           np.where(iris["Target"] == 1, "Versicolor",
                                      "Virginica"))
features = pd.concat([iris["PetalLength"], iris["PetalWidth"],
                      iris["SepalLength"], iris["SepalWidth"]], axis = 1)

from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters = 3, random_state = 29).fit(features)

print(pd.crosstab(index = iris["Species"], columns = kmeans.labels_))

## col_0      0   1   2
## Species
## Setosa      0  50   0

```

```
## Versicolor    2    0  48
## Virginica     36    0  14
```

KMeans

7.2 Spectral Clustering

```
from sklearn.cluster import SpectralClustering
```

```
spectral = SpectralClustering(n_clusters = 3,
                              random_state = 29).fit(features)
```

```
print(pd.crosstab(index = iris["Species"], columns = spectral.labels_))
```

```
## col_0         0    1    2
## Species
## Setosa         0  50    0
## Versicolor    48    0    2
## Virginica     13    0   37
```

SpectralClustering

7.3 Ward Hierarchical Clustering

```
from sklearn.cluster import AgglomerativeClustering
```

```
aggl = AgglomerativeClustering(n_clusters = 3).fit(features)
```

```
print(pd.crosstab(index = iris["Species"], columns = aggl.labels_))
```

```
## col_0         0    1    2
## Species
## Setosa         0  50    0
## Versicolor    49    0    1
## Virginica     15    0   35
```

AgglomerativeClustering

7.4 DBSCAN

```
from sklearn.cluster import DBSCAN
```

```
dbscan = DBSCAN().fit(features)
```

```
print(pd.crosstab(index = iris["Species"], columns = dbscan.labels_))
```

```
## col_0        -1    0    1
## Species
## Setosa         1  49    0
## Versicolor     6    0  44
## Virginica     10    0  40
```

DBSCAN


7.5 Self-organizing map

```
from pyclustering.nnet import som
```

```
sm = som.som(4,4)
```

```
sm.train(features.as_matrix(), 100)
```

```
sm.show_distance_matrix()
```



Output:

[pyclustering](#)

8 Forecasting

8.1 Fit an ARIMA model to a timeseries.

a) Plot the timeseries.

```
# Read in new data set
```


```
air = pd.read_csv('/Users/air.csv')
```

```
air["DATE"] = pd.to_datetime(air["DATE"], infer_datetime_format = True)
```

```
air.index = air["DATE"].values
```

```
plt.plot(air.index, air["AIR"])
```

```
plt.show()
```




Output:

```
to_datetime()
```

b) Fit an ARIMA model and predict 2 years (24 months).

```
import pyflux as pf
```

```
# ar = 12 is necessary to indicate the seasonality of data
model = pf.ARIMA(data = air, ar = 12, ma = 1, integ = 0, target = 'AIR',
                  family = pf.Normal())
x = model.fit("MLE")
model.plot_predict(h = 24)
```



Output:

PyFlux

8.2 Fit a Simple Exponential Smoothing model to a timeseries.

a) Plot the timeseries.

```
# Read in new data set
usecon = pd.read_csv('/Users/usecon.csv')

petrol = usecon["PETROL"]

plt.plot(petrol)
plt.show()
```




Output:

b) Fit a Simple Exponential Smoothing model, predict 2 years (24 months) out and plot predictions.

Currently, there is not a good package in Python to fit a simple exponential smoothing model. The formula for fitting an exponential smoothing model is not difficult, so we can do it by creating our own functions in Python.

The simplest form of exponential smoothing is given by, where $t > 0$:

$$Eq1: s_0 = x_0$$

$$Eq2: s_t = \alpha x_t + (1 - \alpha)s_{t-1}$$

Therefore, we can implement a simple exponential smoothing model as follows:


```
def simple_exp_smoothing(data, alpha, n_preds):
    # Eq1:
    output = [data[0]]
    # Smooth given data plus we want to predict 24 units
    # past the end
    for i in range(1, len(data) + n_preds):
        # Eq2:
        if (i < len(data)):
            output.append(alpha * data[i] + (1 - alpha) * data[i-1])
        else:
            output.append(alpha * output[i-1] + (1 - alpha) * output[i-2])
    return output
```

```

pred = simple_exp_smoothing(petrol, 0.9999, 24)

plt.plot(pd.DataFrame(pred), color = "red")
plt.plot(petrol, color = "blue")
plt.show()

```



Output:

[Basis for code](#)

8.3 Fit a Holt-Winters model to a timeseries.


a) Plot the timeseries.

```
vehicle = usecon["VEHICLE"]
```

```

plt.plot(vehicle)
plt.show()

```



Output:

b) Fit a Holt-Winters additive model, predict 2 years (24 months) out and plot predictions.

Currently, there is not a good package in Python to fit a Holt-Winters additive model. The formula for fitting a Holt-Winters additive model is not difficult, so we can do it by creating our own functions in Python.

The following is an implementation of the Holt-Winters additive model given at [triple exponential smoothing code](#).

```
def initial_trend(series, slen):
    sum = 0.0
    for i in range(slen):
        sum += float(series[i+slen] - series[i]) / slen
    return sum / slen

def initial_seasonal_components(series, slen):
    seasonals = {}
    season_averages = []
    n_seasons = int(len(series)/slen)
    # compute season averages
    for j in range(n_seasons):
        season_averages.append(sum(series[slen*j:slen*j+slen])/float(slen))
    # compute initial values
    for i in range(slen):
        sum_of_vals_over_avg = 0.0
        for j in range(n_seasons):
            sum_of_vals_over_avg += series[slen*j+i]-season_averages[j]
        seasonals[i] = sum_of_vals_over_avg/n_seasons
    return seasonals
```


```

def triple_exponential_smoothing_add(series, slen, alpha, beta, gamma,
                                     n_preds):
    result = []
    seasonals = initial_seasonal_components(series, slen)
    for i in range(len(series)+n_preds):
        if i == 0: # initial values
            smooth = series[0]
            trend = initial_trend(series, slen)
            result.append(series[0])
            continue
        if i >= len(series): # we are forecasting
            m = i - len(series) + 1
            result.append((smooth + m*trend) + seasonals[i%slen])
        else:
            val = series[i]
            last_smooth, smooth = smooth, alpha*(val-seasonals[i%slen])
                                                    + (1-alpha)*(smooth+trend)
            trend = beta * (smooth-last_smooth) + (1-beta)*trend
            seasonals[i%slen] = gamma*(val-smooth) +
                                   (1-gamma)*seasonals[i%slen]
            result.append(smooth+trend+seasonals[i%slen])
    return result

add_preds = triple_exponential_smoothing_add(vehicles, 12, 0.5731265, 0,
                                             0.7230956, 24)

plt.plot(pd.DataFrame(add_preds), color = "red")
plt.plot(vehicles, color = "blue")
plt.show()

```



Output:

9 Model Evaluation & Selection

9.1 Evaluate the accuracy of regression models.

a) Evaluation on training data.

```
train = pd.read_csv('/Users/boston_train.csv')
test = pd.read_csv('/Users/boston_test.csv')

# Random Forest Regression Model
from sklearn.ensemble import RandomForestRegressor
rfMod = RandomForestRegressor(random_state=29)
rfMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Evaluation on training data
predY = rfMod.predict(train.drop(["Target"], axis = 1))

# Determine coefficient of determination score
from sklearn.metrics import r2_score
r2_rf = r2_score(train["Target"], predY)
print("Random forest regression model r^2 score (coefficient of
determination): %f" % r2_rf)

## Random forest regression model r^2 score (coefficient of determination):
0.975233
```

b) Evaluation on testing data.

```
# Random Forest Regression Model (rfMod)

# Evaluation on testing data
predY = rfMod.predict(test.drop(["Target"], axis = 1))

# Determine coefficient of determination score
r2_rf = r2_score(test["Target"], predY)
print("Random forest regression model r^2 score (coefficient of
determination): %f" % r2_rf)

## Random forest regression model r^2 score (coefficient of determination):
0.833687
```

RandomForestRegressor

The sklearn metric `r2_score` is only one option for assessing a regression model. Please go [here](#) for more information about other sklearn regression metrics.

9.2 Evaluate the accuracy of classification models.

a) Evaluation on training data.

```
train = pd.read_csv('/Users/digits_train.csv')
test = pd.read_csv('/Users/digits_test.csv')

# Random Forest Classification Model
from sklearn.ensemble import RandomForestClassifier
rfMod = RandomForestClassifier(random_state=29)
rfMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Evaluation on training data
predY = rfMod.predict(train.drop(["Target"], axis = 1))

# Determine accuracy score
from sklearn.metrics import accuracy_score
accuracy_rf = accuracy_score(train["Target"], predY)
print("Random forest model accuracy: %f" % accuracy_rf)

## Random forest model accuracy: 1.000000
```

b) Evaluation on testing data.

```
# Random Forest Classification Model (rfMod)

# Evaluation on testing data
predY = rfMod.predict(test.drop(["Target"], axis = 1))

# Determine accuracy score
accuracy_rf = accuracy_score(test["Target"], predY)
print("Random forest model accuracy: %f" % accuracy_rf)

## Random forest model accuracy: 0.940741
```

RandomForestClassifier

Note: The sklearn metric `accuracy_score` is only one option for assessing a classification model. Please go [here](#) for more information about other sklearn classification metrics.

9.3 Evaluation with cross validation.

a) KFold

```
# Notice we are using a new data set that need to be read into the
# environment
breastcancer = pd.read_csv('/Users/breastcancer.csv')

from sklearn import model_selection
from sklearn.ensemble import RandomForestClassifier

X = breastcancer.drop(["Target"], axis = 1)
```

```
Y = breastcancer["Target"]

kfold = model_selection.KFold(n_splits = 5, random_state = 29)
rfMod = RandomForestClassifier(random_state = 29)
results = model_selection.cross_val_score(rfMod, X, Y, cv = kfold)

print("Accuracy: %.2f%% +/- %.2f%%" % (results.mean()*100,
                                       results.std()*100))

## Accuracy: 94.38% +/- 2.39%
```

RandomForestClassifier | KFold

b) ShuffleSplit

```
shuffle = model_selection.ShuffleSplit(n_splits = 5, random_state = 29)
rfMod = RandomForestClassifier(random_state = 29)
results = model_selection.cross_val_score(rfMod, X, Y, cv = shuffle)

print("Accuracy: %.2f%% +/- %.2f%%" % (results.mean()*100,
                                       results.std()*100))

## Accuracy: 95.09% +/- 0.70%
```

RandomForestClassifier | ShuffleSplit

10 Text Analytics

11 Deep Learning

Appendix

1 Built-in Python Data Types

- Boolean

Numeric types

- int
- long
- float
- complex

Sequences

- str
- bytes
- byte array
- list
- tuple

Sets

- set
- frozen set

Mapping:

- dictionary

2 Python Plotting Packages

Bokeh

A Python package which is useful for interactive visualizations and is optimized for web browser presentations.

PyPlot

A Python package which is useful for data plotting and visualization.

Seaborn

A Python package which is useful for data plotting and visualization. In particular, Seaborn includes tools for drawing attractive statistical graphics.

3 Python packages used in this tutorial

pandas

Working with data structures and performing data analysis

NumPy

Scientific and mathematical computing

re

Regular expressions

Decimal

Tools for decimal floating point arithmetic

sklearn

scikit-learn, or more commonly known as sklearn, is useful for basic and advanced data mining, machine learning, and data analysis. sklearn includes tools for classification, regression, clustering, dimensionality reduction, model selection, and data pre-processing.

statsmodels.api

Tools for the estimation of many different statistical models

xgboost

Extreme gradient boosting models

pyclustering

Tools for clustering input data

PyFlux

Tools for time series analysis and prediction

Alphabetical Index

Array

A **NumPy** array is a data type in which the elements of the array are all of the same type. Please see the following example of array creation and access:

```
import numpy as np
my_array = np.array([1, 2, 3, 4])
print(my_array)

## [1 2 3 4]

print(my_array[3])

## 4
```

Bytes & Byte arrays

A **byte** is a sequence of integers which is immutable, whereas a **byte array** is its mutable counterpart.

Data Frame

A two-dimensional tabular structure with labeled axes (rows and columns), where data observations are represented by rows and data variables are represented by columns.

Dictionary

An associative array which is indexed by keys which map to values. Therefore, a dictionary is an unordered set of key:value pairs where each key is unique. Please see the following example of dictionary creation and access:

```
import pandas as pd
student = pd.read_csv('/Users/class.csv')
for_dict = pd.concat([student["Name"], student["Age"]], axis = 1)
class_dict = for_dict.set_index('Name').T.to_dict('list')
print(class_dict.get('James'))

## [12]
```

List

A sequence of comma-separated objects that need not be of the same type. Please see the following example of list creation and access:

```
list1 = ['item1', 102]
print(list1)

## ['item1', 102]

print(list1[1])

## 102
```

Python also has what are known as "Tuples", which are immutable lists created in the same way as lists, except with paranthesis instead of brackets.

Series

A one-dimensional data frame. Please see the following example of Series creation and access:

```
import pandas as pd
my_array = pd.Series([1, 3, 5, 9])
print(my_array)

## 0    1
## 1    3
## 2    5
## 3    9
## dtype: int64

print(my_array[1])

## 3
```

Sets & Frozen Sets

A set is an unordered collection of immutable objects. The difference between a [set](#) and a [frozen set](#) is that the former is mutable, while the latter is immutable. Please see the following example of set and frozen set creation and access:

```
s = set(["1", "2", "3"])
# s is a set, which means you can add or delete elements from s
print(s)

## {'2', '1', '3'}

s.add("4")
print(s)

## {'3', '4', '1', '2'}

fs = frozenset(["1", "2", "3"])
# fs is a frozenset, which means you cannot add or delete elements from fs
print(fs)

## frozenset({'2', '1', '3'})
```

str

A list of characters, though characters are not a type in Python, but rather a string of length 1. Strings are indexable like arrays. Please see the following example of String creation and access:

```
s = 'My first string!'
print(s)

## My first string!

print(s[5])

## r
```

For more information on Python packages and functions, along with helpful examples, please see [Python](#).