

# Python Tutorial

First, you need to import several important Python packages for data manipulation and scientific computing. The [pandas](#) package is helpful for data manipulation and the [NumPy](#) package is helpful for scientific computing.

```
import pandas as pd
import numpy as np
```

In Python, comments are indicated in code with a "#" character, and arrays and matrices are zero-indexed.

## 1 Reading in Data and Basic Statistical Functions

### 1.1 Read in the data.

The following demonstrate importing data in Python given 3 different data formats. The pandas package is able to read all 3 formats, as well as many others, using [Python IO tools](#).

#### a) Read the data in as a .csv file.

```
student = pd.read_csv('/Users/class.csv')
```

#### b) Read the data in as a .xls file.

```
# Notice you must specify the file location, as well as the name of the sheet
# of the .xls file you want to import
student_xls = pd.read_excel(open('/Users/class.xls', 'rb'),
                             sheetname='class')
```

#### c) Read the data in as a .json file.

```
student_json = pd.read_json('/Users/class.json')
```

### 1.2 Find the dimensions of the data set.

The dimensions of a [DataFrame](#) in Python are known as an attribute of the object. Therefore, you can state the data name followed by ".shape" to return the dimensions of the data.

```
print(student.shape)
## (19, 5)
```

### 1.3 Find basic information about the data set.

Information about a [DataFrame](#) is available by calling the ".info()" function on the data.

```
# Notice that student is a DataFrame object
print(student.info())

## <class 'pandas.core.frame.DataFrame'>
## RangeIndex: 19 entries, 0 to 18
## Data columns (total 5 columns):
## Name      19 non-null object
## Sex       19 non-null object
## Age       19 non-null int64
## Height    19 non-null float64
## Weight    19 non-null float64
## dtypes: float64(2), int64(1), object(2)
## memory usage: 840.0+ bytes
## None
```

## 1.4 Look at the first 5 observations.

The first 5 observations of a [DataFrame](#) are available by calling the ".head()" function on the data. By default, head() returns 5 observations. To return the first  $n$  observations, pass the integer  $n$  into the function. The tail() function is analogous and returns the last observations.

```
print(student.head())

##      Name Sex  Age  Height  Weight
## 0  Alfred   M   14    69.0   112.5
## 1   Alice   F   13    56.5    84.0
## 2 Barbara   F   13    65.3    98.0
## 3   Carol   F   14    62.8   102.5
## 4   Henry   M   14    63.5   102.5
```

## 1.5 Calculate mean of numeric variables.

```
# By default, the mean() function returns the mean of numeric variables of
# the data only
print(student.mean())

## Age      13.315789
## Height   62.336842
## Weight   100.026316
## dtype: float64
```

## 1.6 Compute summary statistics of the data set.

Summary statistics of a [DataFrame](#) are available by calling the ".describe()" function on the data.

```
print(student.describe())

##      Age      Height      Weight
## count  19.000000  19.000000  19.000000
```

```
## mean    13.315789  62.336842  100.026316
## std      1.492672   5.127075   22.773933
## min     11.000000  51.300000   50.500000
## 25%     12.000000  58.250000   84.250000
## 50%     13.000000  62.800000   99.500000
## 75%     14.500000  65.900000  112.250000
## max     16.000000  72.000000  150.000000
```

## 1.7 Descriptive statistics functions applied to variables of the data set.

*# Notice the subsetting of student with [] and the name of the variable in # quotes*

```
print(student["Weight"].std())
## 22.773933493879046

print(student["Weight"].sum())
## 1900.5

print(student["Weight"].count())
## 19

print(student["Weight"].max())
## 150.0

print(student["Weight"].min())
## 50.5

print(student["Weight"].median())
## 99.5
```

## 1.8 Produce a one-way table to describe the frequency of a variable.

### a) Produce a one-way table of a discrete variable.

*# columns = "count" indicates to make the descriptive portion of the table  
# the counts of each level of the index variable*

```
print(pd.crosstab(index=student["Age"], columns="count"))
## col_0  count
## Age
## 11        2
## 12        5
## 13        3
## 14        4
## 15        4
## 16        1
```

## b) Produce a one-way table of a categorical variable.

```
print(pd.crosstab(index=student["Sex"], columns="count"))

## col_0  count
## Sex
## F          9
## M         10
```

## 1.9 Produce a two-way table to describe the frequency of two categorical or discrete variables.

*# Notice the specification of a variable for the columns argument, instead  
# of "count"*

```
## Sex  F  M
## Age
## 11   1  1
## 12   2  3
## 13   2  1
## 14   2  2
## 15   2  2
## 16   0  1
```

`crosstab()`

## 1.10 Select a subset of the data that meets a certain criterion.

```
females = student.query('Sex == "F"')
print(females.head())

##      Name Sex  Age  Height  Weight
## 1  Alice   F   13   56.5    84.0
## 2 Barbara   F   13   65.3    98.0
## 3  Carol   F   14   62.8   102.5
## 6   Jane   F   12   59.8    84.5
## 7  Janet   F   15   62.5   112.5
```

`query()`

## 1.11 Determine the correlation between two continuous variables.

```
height_weight = pd.concat([student["Height"], student["Weight"]], axis = 1)
print(height_weight.corr(method = "pearson"))

##      Height  Weight
## Height  1.000000  0.877785
## Weight  0.877785  1.000000
```

`corr()`

---

## 2 Basic Graphing and Plotting Functions


The Matplotlib [PyPlot](#) package is a standard Python package to use for plotting. For more information on other Python plotting packages, please see the Appendix Section 2.

```
import matplotlib.pyplot as plt
```

### 2.1 Visualize a single continuous variable by producing a histogram.

```
# Notice how the bin endpoints are set so the histogram is the same  
# as that produced by SAS and R  
# Also notice the labeling of the axes
```


```
plt.hist(student["Weight"], bins=[40,60,80,100,120,140,160])  
plt.xlabel('Weight')  
plt.ylabel('Frequency')  
plt.show()
```



Output:

### 2.2 Visualize a single continuous variable by producing a boxplot.

```
# showmeans=True tells Python to plot the mean of the variable on the boxplot  
plt.boxplot(student["Weight"], showmeans=True)  
# prevents Python from printing a "1" at the bottom of the boxplot  
plt.xticks([]) # prevents Python from printing a "1" at the bottom of the boxplot  
plt.ylabel('Weight')  
plt.show()
```



Output:

### 2.3 Visualize two continuous variables by producing a scatterplot.


*# Notice here you specify the x variable first followed by the y variable*

```
plt.scatter(student["Height"], student["Weight"])
```

```
plt.xlabel("Height")
```

```
plt.ylabel("Weight")
```

```
plt.show()
```




Output:

## 2.4 Visualize a relationship between two continuous variables by producing a scatterplot and a plotted line of best fit.

```
x = student["Height"]
y = student["Weight"]

# np.polyfit() models Weight as a function of Height and returns the
# parameters
m, b = np.polyfit(x, y, 1)
plt.scatter(x, y)

# plt.text() prints the equation of the line of best fit, with the first two
# arguments specifying the x and y locations of the text, respectively
# %f indicates to print a floating point number, that is specified following
# the string
plt.text(51, 140, "Line: y = %f x + %f" % (m,b))
plt.plot(x, m*x + b)
plt.xlabel("Height")
plt.ylabel("Weight")
plt.show()
```



Output:

NumPy polyfit()

## 2.5 Visualize a categorical variable by producing a bar chart.


```
# Get the counts of Sex
counts = pd.crosstab(index=student["Sex"], columns="count")
```

```

# len() returns the number of categories of Sex (2)
# np.arange() creates a vector of the specified length
num = np.arange(len(counts))
plt.bar(num, counts["count"], align='center', alpha=0.5)

# Set the xticks to be the indices of counts
plt.xticks(num, counts.index)
plt.xlabel("Sex")
plt.ylabel("Frequency")
plt.show()

```



Output:

NumPy arange()

## 2.6 Visualize a continuous variable, grouped by a categorical variable, by producing side-by-side boxplots.

### a) Simple side-by-side boxplot without color.

```


# Subset data set to return only female weights, and then only male weights
Weight_F = np.array(student.query('Sex == "F"')['Weight'])
Weight_M = np.array(student.query('Sex == "M"')['Weight'])
Weights = [Weight_F, Weight_M]

# PyPlot automatically plots the two weights side-by-side since Weights
# is a 2D array
plt.boxplot(Weights, showmeans=True, labels=('F', 'M'))
plt.xlabel('Sex')

```



```
plt.ylabel('Weight')
plt.show()
```




Output:

#### b) More advanced side-by-side boxplot with color.

```
import seaborn as sns
sns.boxplot(x="Sex", y="Weight", hue="Sex", data = student, showmeans=True)
plt.show()
```

[seaborn boxplot](#)

[seaborn](#)



Output:

### 3 Basic Data Wrangling and Manipulation

#### 3.1 Create a new variable in a data set as a function of existing variables in the data set.

*# Notice here how you can create the BMI column in the data set just by naming it*

```
student["BMI"] = student["Weight"] / student["Height"]**2 * 703
print(student.head())
```

##	Name	Sex	Age	Height	Weight	BMI
## 0	Alfred	M	14	69.0	112.5	16.611531
## 1	Alice	F	13	56.5	84.0	18.498551
## 2	Barbara	F	13	65.3	98.0	16.156788
## 3	Carol	F	14	62.8	102.5	18.270898
## 4	Henry	M	14	63.5	102.5	17.870296

#### 3.2 Create a new variable in a data set using if/else logic of existing variables in the data set.

*# Notice the use of the np.where() function for a single condition*

```
student["BMI Class"] = np.where(student["BMI"] < 19.0, "Underweight",
```

```
"Healthy")
print(student.head())
NA
```

	Name	Sex	Age	Height	Weight	BMI	BMI Class
## 0	Alfred	M	14	69.0	112.5	16.611531	Underweight
## 1	Alice	F	13	56.5	84.0	18.498551	Underweight
## 2	Barbara	F	13	65.3	98.0	16.156788	Underweight
## 3	Carol	F	14	62.8	102.5	18.270898	Underweight
## 4	Henry	M	14	63.5	102.5	17.870296	Underweight

NumPy where()

### 3.3 Create new variables in a data set using mathematical functions applied to existing variables in the data set.

Using the log() function, the exp() function, the sqrt() function, and the abs() function.

```
student["LogWeight"] = np.log(student["Weight"])
student["ExpAge"] = np.exp(student["Age"])
student["SqrtHeight"] = np.sqrt(student["Height"])
student["BMI Neg"] = np.where(student["BMI"] < 19.0, -student["BMI"],
student["BMI"])
student["BMI Pos"] = np.abs(student["BMI Neg"])

# Create a boolean variable
student["BMI Check"] = (student["BMI Pos"] == student["BMI"])
print(student.head())
```

	Name	Sex	Age	Height	Weight	BMI	BMI Class	LogWeight	\
## 0	Alfred	M	14	69.0	112.5	16.611531	Underweight	4.722953	
## 1	Alice	F	13	56.5	84.0	18.498551	Underweight	4.430817	
## 2	Barbara	F	13	65.3	98.0	16.156788	Underweight	4.584967	
## 3	Carol	F	14	62.8	102.5	18.270898	Underweight	4.629863	
## 4	Henry	M	14	63.5	102.5	17.870296	Underweight	4.629863	

	ExpAge	SqrtHeight	BMI Neg	BMI Pos	BMI Check
## 0	1.202604e+06	8.306624	-16.611531	16.611531	True
## 1	4.424134e+05	7.516648	-18.498551	18.498551	True
## 2	4.424134e+05	8.080842	-16.156788	16.156788	True
## 3	1.202604e+06	7.924645	-18.270898	18.270898	True
## 4	1.202604e+06	7.968689	-17.870296	17.870296	True

### 3.4 Drop variables from a data set.

```
# axis = 1 indicates to drop columns instead of rows
student = student.drop(["LogWeight", "ExpAge", "SqrtHeight", "BMI Neg",
"BMI Pos", "BMI Check"], axis = 1)
print(student.head())
```

	Name	Sex	Age	Height	Weight	BMI	BMI Class
## 0	Alfred	M	14	69.0	112.5	16.611531	Underweight
## 1	Alice	F	13	56.5	84.0	18.498551	Underweight
## 2	Barbara	F	13	65.3	98.0	16.156788	Underweight
## 3	Carol	F	14	62.8	102.5	18.270898	Underweight
## 4	Henry	M	14	63.5	102.5	17.870296	Underweight

### 3.5 Sort a data set by a variable.

#### a) Sort data set by a continuous variable.

*# Notice the kind="mergesort" which indicates to use a stable sorting algorithm*

```
student = student.sort_values(by="Age", kind = "mergesort")
print(student.head())
```

	Name	Sex	Age	Height	Weight	BMI	BMI Class
## 10	Joyce	F	11	51.3	50.5	13.490001	Underweight
## 17	Thomas	M	11	57.5	85.0	18.073346	Underweight
## 5	James	M	12	57.3	83.0	17.771504	Underweight
## 6	Jane	F	12	59.8	84.5	16.611531	Underweight
## 9	John	M	12	59.0	99.5	20.094369	Healthy

#### b) Sort data set by a categorical variable.

```
student = student.sort_values(by="Sex", kind = "mergesort")
# Notice that the data is now sorted first by Sex and then within Sex by Age
print(student.head())
```

	Name	Sex	Age	Height	Weight	BMI	BMI Class
## 10	Joyce	F	11	51.3	50.5	13.490001	Underweight
## 6	Jane	F	12	59.8	84.5	16.611531	Underweight
## 12	Louise	F	12	56.3	77.0	17.077695	Underweight
## 1	Alice	F	13	56.5	84.0	18.498551	Underweight
## 2	Barbara	F	13	65.3	98.0	16.156788	Underweight

### 3.6 Compute descriptive statistics of continuous variables, grouped by a categorical variable.

```
print(student.groupby(by = "Sex").mean())
```

	Age	Height	Weight	BMI
## Sex				
## F	13.222222	60.588889	90.111111	17.051039
## M	13.400000	63.910000	108.950000	18.594243

### 3.7 Add a new row to the bottom of a data set.

*# Look at the tail of the data currently*

```
print(student.tail())
```

	Name	Sex	Age	Height	Weight	BMI	BMI Class
## 0	Alfred	M	14	69.0	112.5	16.611531	Underweight

```
## 4      Henry    M    14     63.5    102.5    17.870296    Underweight
## 16     Ronald    M    15     67.0    133.0    20.828470         Healthy
## 18    William    M    15     66.5    112.0    17.804511    Underweight
## 14     Philip    M    16     72.0    150.0    20.341435         Healthy

student = student.append({'Name':'Jane', 'Sex':'F', 'Age':14, 'Height':56.3,
                          'Weight':77.0, 'BMI':17.077695,
                          'BMI Class': 'Underweight'},
                          ignore_index=True)
```

*# Notice the change in the indices because of the ignore\_index=True option  
# which allows for a Series, or one-dimensional DataFrame, to be appended  
# to an existing DataFrame*

```
##      Name Sex  Age  Height  Weight      BMI  BMI Class
## 15   Henry  M   14    63.5   102.5  17.870296 Underweight
## 16   Ronald  M   15    67.0   133.0  20.828470     Healthy
## 17   William  M   15    66.5   112.0  17.804511 Underweight
## 18   Philip  M   16    72.0   150.0  20.341435     Healthy
## 19    Jane   F   14    56.3    77.0  17.077695 Underweight
```

### 3.8 Create a user defined function and apply it to a variable in the data set to create a new variable in the data set.

```
def toKG(lb):
    return (0.45359237 * lb)

student["Weight KG"] = student["Weight"].apply(toKG)
print(student.head())
```

```
##      Name Sex  Age  Height  Weight      BMI  BMI Class  Weight KG
## 0    Joyce  F   11    51.3    50.5  13.490001 Underweight  22.906415
## 1     Jane  F   12    59.8    84.5  16.611531 Underweight  38.328555
## 2   Louise  F   12    56.3    77.0  17.077695 Underweight  34.926612
## 3    Alice  F   13    56.5    84.0  18.498551 Underweight  38.101759
## 4  Barbara  F   13    65.3    98.0  16.156788 Underweight  44.452052
```

`apply()`

### 3.9 Caste a Data Frame to a different object type.

```
student_num = pd.concat([student["Age"], student["Height"]],

##      Age  Height  Weight
## 0  14.0    69.0   112.5
## 1  13.0    56.5    84.0
## 2  13.0    65.3    98.0
## 3  14.0    62.8   102.5
## 4  14.0    63.5   102.5
```

`astype()`

---

## 4 More Advanced Data Wrangling

### 4.1 Drop observations with missing information.

```
# Notice the use of the fish data set because it has some missing
# observations
fish = pd.read_csv('/Users/fish.csv')

# First sort by Weight, requesting those with NA for Weight first
fish = fish.sort_values(by='Weight', kind='mergesort', na_position='first')
print(fish.head())
```

	Species	Weight	Length1	Length2	Length3	Height	Width
## 13	Bream	NaN	29.5	32.0	37.3	13.9129	5.0728
## 40	Roach	0.0	19.0	20.5	22.8	6.4752	3.3516
## 72	Perch	5.9	7.5	8.4	8.8	2.1120	1.4080
## 145	Smelt	6.7	9.3	9.8	10.8	1.7388	1.0476
## 147	Smelt	7.0	10.1	10.6	11.6	1.7284	1.1484

--

```
new_fish = fish.dropna()
print(new_fish.head())
```

	Species	Weight	Length1	Length2	Length3	Height	Width
## 40	Roach	0.0	19.0	20.5	22.8	6.4752	3.3516
## 72	Perch	5.9	7.5	8.4	8.8	2.1120	1.4080
## 145	Smelt	6.7	9.3	9.8	10.8	1.7388	1.0476
## 147	Smelt	7.0	10.1	10.6	11.6	1.7284	1.1484
## 146	Smelt	7.5	10.0	10.5	11.6	1.9720	1.1600

`dropna()`

### 4.2 Merge two data sets together on a common variable.

#### a) First, select specific columns of a data set to create two smaller data sets.

```
# Notice the use of the student data set again, however we want to reload it
# without the changes we've made previously
student = pd.read_csv('/Users/class.csv')
student1 = pd.concat([student["Name"], student["Sex"], student["Age"]],
                    axis = 1)
print(student1.head())
NA
```

	Name	Sex	Age
## 0	Alfred	M	14
## 1	Alice	F	13
## 2	Barbara	F	13

```
## 3    Carol    F    14
## 4    Henry    M    14

--

student2 = pd.concat([student["Name"], student["Height"], student["Weight"]],
                     axis = 1)
print(student2.head())
```

	Name	Height	Weight
## 0	Alfred	69.0	112.5
## 1	Alice	56.5	84.0
## 2	Barbara	65.3	98.0
## 3	Carol	62.8	102.5
## 4	Henry	63.5	102.5

**b) Second, we want to merge the two smaller data sets on the common variable.**

```
new = pd.merge(student1, student2, on="Name")
print(new.head())
```

	Name	Sex	Age	Height	Weight
## 0	Alfred	M	14	69.0	112.5
## 1	Alice	F	13	56.5	84.0
## 2	Barbara	F	13	65.3	98.0
## 3	Carol	F	14	62.8	102.5
## 4	Henry	M	14	63.5	102.5

**c) Finally, we want to check to see if the merged data set is the same as the original data set.**

```
print(student.equals(new))
```

```
## True
```

`merge()`

## 4.3 Merge two data sets together by index number only.

**a) First, select specific columns of a data set to create two smaller data sets.**

```
newstudent1 = pd.concat([student["Name"], student["Sex"], student["Age"]],
                       axis = 1)
print(newstudent1.head())
```

	Name	Sex	Age
## 0	Alfred	M	14
## 1	Alice	F	13
## 2	Barbara	F	13
## 3	Carol	F	14
## 4	Henry	M	14

##	Height	Weight
## 0	69.0	112.5
## 1	56.5	84.0
## 2	65.3	98.0
## 3	62.8	102.5
## 4	63.5	102.5

```
new2 = newstudent1.join(newstudent2)
print(new2.head())
```

##	Name	Sex	Age	Height	Weight
## 0	Alfred	M	14	69.0	112.5
## 1	Alice	F	13	56.5	84.0
## 2	Barbara	F	13	65.3	98.0
## 3	Carol	F	14	62.8	102.5
## 4	Henry	M	14	63.5	102.5

```
print(student.equals(new2))
```

```
## True
```

## join()

```
# Notice we are using a new data set that needs to be read into the
# environment
```

```
# The following code is used to remove the ',' and '$' characters from
# the ACTUAL column so that the values can be summed
```

```
price["REVENUE"] = price["ACTUAL"].apply(trim_money)
table = pd.pivot_table(price, index=["COUNTRY", "STATE", "PRODTYPE",
                                     "PRODUCT"], values="REVENUE",
aggfunc=np.sum)
print(table.head())
```

##	COUNTRY	STATE	PRODTYPE	PRODUCT	
##	Canada	British Columbia	FURNITURE	BED	197706.6
##				SOFA	216282.6



```
##                OFFICE    CHAIR    200905.2
##                DESK      186262.2
##    Ontario    FURNITURE    BED      194493.6
## Name: REVENUE, dtype: float64
```

sub

Decimal

[pivot\\_table\(\)](#) for more information on the pandas [pivot\\_table\(\)](#) function

[pivot\(\)](#)

---

## 5 Regression & Modeling

The following sections focus on the Python [sklearn](#) package.

### 5.1 Pre-process a data set using principal component analysis.

```
from sklearn.decomposition import PCA

# Notice we are using a new data set that needs to be read into the
# environment
iris = pd.read_csv('/Users/iris.csv')
features = iris.drop(["Target"], axis = 1)

pca = PCA(n_components = 4)
pca = pca.fit(features)
print(np.transpose(pca.components_))

## [[ 0.36158968  0.65653988 -0.58099728  0.31725455]
##   [-0.08226889  0.72971237  0.59641809 -0.32409435]
##   [ 0.85657211 -0.1757674   0.07252408 -0.47971899]
##   [ 0.35884393 -0.07470647  0.54906091  0.75112056]]
```

PCA

### 5.2 Split data into training and testing data and export as a .csv file.

```
from sklearn.model_selection import train_test_split

target = iris["Target"]

# The following code splits the iris data set into 70% train and 30% test
X_train, X_test, Y_train, Y_test = train_test_split(features, target,
                                                    test_size = 0.3,
                                                    random_state = 29)

train_x = pd.DataFrame(X_train)
train_y = pd.DataFrame(Y_train)
```

```

test_x = pd.DataFrame(X_test)
test_y = pd.DataFrame(Y_test)

train = pd.concat([train_x, train_y], axis = 1)
test = pd.concat([test_x, test_y], axis = 1)

train.to_csv('/Users/iris_train.csv', index = False)
test.to_csv('/Users/iris_test.csv', index = False)

```

[train\\_test\\_split](#)

### 5.3 Fit a logistic regression model.

```

# Notice we are using a new data set that needs to be read into the
# environment
tips = pd.read_csv('/Users/tips.csv')

# The following code is used to determine if the individual left more
# than a 15% tip
tips["fifteen"] = 0.15 * tips["total_bill"]
tips["greater15"] = np.where(tips["tip"] > tips["fifteen"], 1, 0)

import statsmodels.api as sm

# Notice the syntax of greater15 as a function of total_bill
res = sm.formula.glm("greater15 ~ total_bill", family=sm.families.Binomial(),
                    data=tips).fit()
print(res.summary())

##                               Generalized Linear Model Regression Results
##
=====
=
## Dep. Variable:                greater15    No. Observations:
244
## Model:                        GLM          Df Residuals:
242
## Model Family:                 Binomial     Df Model:
1
## Link Function:                logit        Scale:
1.0
## Method:                       IRLS        Log-Likelihood:
156.87
## Date:                         Wed, 07 Jun 2017    Deviance:
313.74
## Time:                         16:27:29          Pearson chi2:
247.
## No. Iterations:                6
##
=====

```

```

=
##              coef      std err          z      P>|z|      [95.0% Conf.
Int.]
## -----
----
## Intercept      1.6477      0.355      4.646      0.000      0.953
2.343
## total_bill     -0.0725      0.017     -4.319      0.000     -0.105   -
0.040
##
=====
=

```

A logistic regression model can be implemented using [sklearn](#), however [statsmodels.api](#) provides a helpful summary about the model, so it is preferable for this example.

## 5.4 Fit a linear regression model on training data and assess against testing data.

```

# Notice we are using new data sets that need to be read into the environment
train = pd.read_csv('/Users/tips_train.csv')
test = pd.read_csv('/Users/tips_test.csv')

# Fit a linear regression model of tip by total_bill on the training data
from sklearn import linear_model
regr = linear_model.LinearRegression()
# If your data has one feature, you need to reshape the 1D array
model = regr.fit(train["total_bill"].reshape(-1,1), train["tip"])

# Predict the tip based on the total_bill given in the testing data
prediction = pd.DataFrame()
prediction["tip_hat"] = regr.predict(test["total_bill"].reshape(-1,1))

# Compute the squared difference between predicted tip and actual tip
prediction["diff"] = (prediction["tip_hat"] - test["tip"])**2

# Compute the mean of the squared differences (mean squared error)
# as an assessment of the model
mean_sq_error = np.mean(prediction["diff"])
print(mean_sq_error)

## 1.08759363430702

```

[LinearRegression](#)

## 5.5 Fit a decision tree model on training data and assess against testing data.

a) Build a model, assess the model against the training data, plot the tree, and determine variable importance.

```
# Notice we are using new data sets that need to be read into the environment
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

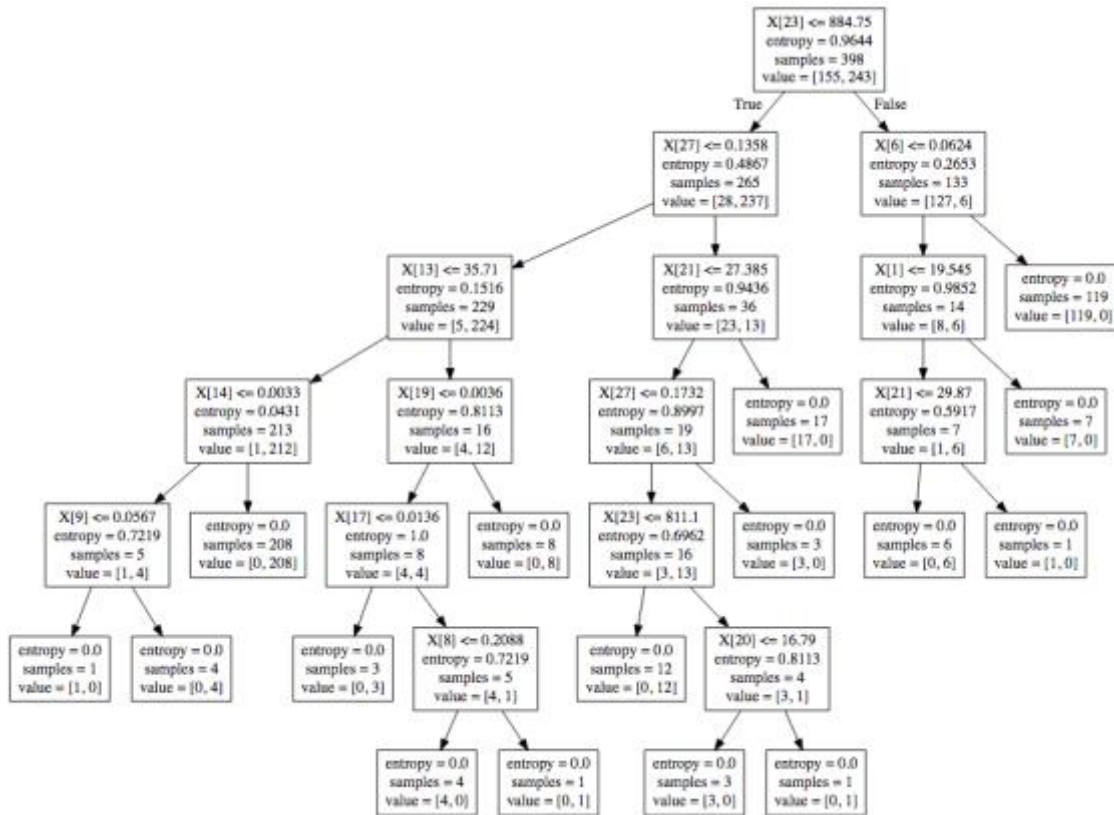
from sklearn import tree

# random_state is used to specify a seed for a random integer so that the
# results are reproducible
clf = tree.DecisionTreeClassifier(criterion='entropy', random_state=29)
clf = clf.fit(train.drop(["Target"], axis = 1), train["Target"])

# Prediction on training data
scored = pd.DataFrame(clf.predict(train.drop(["Target"], axis = 1)))
scored["Target"] = train["Target"]

# Determine how many were correctly classified
scored["correct"] = (scored["Target"] == scored[0])

## col_0    count
## correct
## True      398
```



Output:

```
# Determine variable importance
var_import = clf.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0:'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                     ascending = False)

print(var_import.head())

##      Importance
## 23      0.592658
## 27      0.172561
## 6       0.055977
## 21      0.054751
## 13      0.032737
```

## b) Assess the model against the testing data.

```
# Prediction on testing data
scored = pd.DataFrame(clf.predict(test.drop(["Target"], axis = 1)))
scored["Target"] = test["Target"]

# Determine how many were correctly classified
scored["correct"] = (scored["Target"] == scored[0])
print(pd.crosstab(index=scored["correct"], columns="count"))
```

```
## col_0    count
## correct
## False      9
## True      162
```

## DecisionTreeClassifier

### 5.6 Fit a random forest classification model on training data and assess against testing data.

#### a) Build a model, assess the model against the training data, and determine variable importance.

```
# Notice we are using new data sets that need to be read into the environment
train = pd.read_csv('/Users/iris_train.csv')
test = pd.read_csv('/Users/iris_test.csv')
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
clf = RandomForestClassifier(random_state=29)
clf = clf.fit(train.drop(["Target"], axis = 1), train["Target"])
```

```
# Prediction on training data
scored = pd.DataFrame(clf.predict(train.drop(["Target"], axis = 1)))
scored["Target"] = train["Target"]
```

```
# Determine how many were correctly classified
scored["correct"] = (scored["Target"] == scored[0])
print(pd.crosstab(index=scored["correct"], columns="count"))
```

```
## col_0    count
## correct
## True      105
```

```
--
```

```
# Determine variable importance
var_import = clf.feature_importances_
var_import = pd.DataFrame(var_import)
var_import = var_import.rename(columns = {0:'Importance'})
var_import = var_import.sort_values(by="Importance", kind = "mergesort",
                                     ascending = False)

print(var_import.head())
```

```
##      Importance
## 2      0.515197
## 3      0.388860
## 0      0.080865
## 1      0.015078
```

## b) Assess the model against the testing data.

```
# Prediction on testing data
scored = pd.DataFrame(clf.predict(test.drop(["Target"], axis = 1)))
scored["Target"] = test["Target"]

# Determine how many were correctly classified
scored["correct"] = (scored["Target"] == scored[0])
print(pd.crosstab(index=scored["correct"], columns="count"))

## col_0    count
## correct
## False      3
## True      42
```

RandomForestClassifier

## 5.7 Fit a random forest regression model on training data and assess against testing data.

### a) Build a model and assess the model against the training data.

```
# Notice we are re-using data sets but it is good to re-read the original
# version back into the environment
train = pd.read_csv('/Users/tips_train.csv')
test = pd.read_csv('/Users/tips_test.csv')

from sklearn.ensemble import RandomForestRegressor

clf = RandomForestRegressor(random_state=29)
clf = clf.fit(train.drop(["tip"], axis = 1), train["tip"])

# Prediction on training data
scored = pd.DataFrame(clf.predict(train.drop(["tip"], axis = 1)))
scored["Target"] = train["tip"]

# Determine mean squared error
scored["diff"] = (scored["Target"] - scored[0])**2
print(scored["diff"].mean())

## 0.2177569846153845
```

### b) Assess the model against the testing data.

```
# Prediction on testing data
scored = pd.DataFrame(clf.predict(test.drop(["tip"], axis = 1)))
scored["Target"] = test["tip"]

# Determine mean squared error
scored["diff"] = (scored["Target"] - scored[0])**2
print(scored["diff"].mean())
```

```
## 1.182229489795918
```

RandomForestRegressor

## 5.8 Fit a gradient boosting model on training data and assess against testing data.

### a) Build a model and assess the model against the training data.

*# Notice we are re-using data sets but it is good to re-read the original version back into the environment*

```
train = pd.read_csv('/Users/breastcancer_train.csv')
```

```
test = pd.read_csv('/Users/breastcancer_test.csv')
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

*# n\_estimators = total number of trees to fit which is analogous to the number of iterations*

*# learning\_rate = shrinkage or step-size reduction, whereas a lower*

*# learning rate requires more iterations*

*# min\_samples\_leaf = minimum number of observations in the trees*

*# terminal nodes*

```
clf = GradientBoostingClassifier(random_state = 29, learning_rate = .01,  
min_samples_leaf = 20, n_estimators = 2500)
```

```
clf = clf.fit(train.drop(["Target"], axis = 1), train["Target"])
```

*# Prediction on training data*

```
scores = pd.DataFrame(clf.predict(train.drop(["Target"], axis = 1)))
```

```
scores["Target"] = train["Target"]
```

*# Determine how many were correctly classified*

```
scores["correct"] = (scores["Target"] == scores[0])
```

```
print(pd.crosstab(index = scores["correct"], columns = "count"))
```

```
## col_0    count
```

```
## correct
```

```
## True      398
```

### b) Assess the model against the testing data.

*# Prediction on testing data*

```
scores = pd.DataFrame(clf.predict(test.drop(["Target"], axis = 1)))
```

```
scores["Target"] = test["Target"]
```

*# Determine how many were correctly classified*

```
scores["correct"] = (scores["Target"] == scores[0])
```

```
## col_0    count
```

```
## correct
```



```
## False      4
## True       167
```

GradientBoostingClassifier

## 5.9 Fit a support vector classification model.

### a) Build a model and assess the model against the training data.

```
# Notice we are re-using data sets but it is good to re-read the original
# version back into the environment
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

# First we need to scale the data
from sklearn.preprocessing import StandardScaler

train_features = train.drop(["Target"], axis = 1)
scaler = StandardScaler().fit(np.array(train_features))
train_scaled = scaler.transform(np.array(train_features))
train_scaled = pd.DataFrame(train_scaled)
train_scaled["Target"] = train["Target"]

test_features = test.drop(["Target"], axis = 1)
scaler = StandardScaler().fit(np.array(test_features))
test_scaled = scaler.transform(np.array(test_features))
test_scaled = pd.DataFrame(test_scaled)
test_scaled["Target"] = test["Target"]

# Fit a support vector classification model
from sklearn.svm import SVC
clf = SVC(random_state = 29, kernel = 'linear')
clf = clf.fit(train_scaled.drop(["Target"], axis = 1),
train_scaled["Target"])

# Evaluation on training data
predictions = pd.DataFrame()
predictions["predY"] = clf.predict(train_scaled.drop(["Target"], axis = 1))

# Determine how many were correctly classified
predictions["actual"] = train_scaled["Target"]
predictions["correct"] = (predictions["actual"] == predictions["predY"])
print(pd.crosstab(index = predictions["correct"], columns = "count"))

## col_0    count
## correct
## False      6
## True     392
```

## b) Assess the model against the testing data.

```
# Evaluation on testing data
predictions = pd.DataFrame()
predictions["predY"] = clf.predict(test_scaled.drop(["Target"], axis = 1))

# Determine how many were correctly classified
predictions["actual"] = test_scaled["Target"]
predictions["correct"] = (predictions["actual"] == predictions["predY"])
print(pd.crosstab(index = predictions["correct"], columns = "count"))

## col_0    count
## correct
## False      7
## True     164
```

## 5.10 Fit a support vector regression model.


### a) Generate random data based on a sine curve.

```
# Generate the time variable
t = np.linspace(start = 0, stop = 0.5*np.pi, num = 100)

# Generate the sine curve with uniform noise
y1 = 5*np.sin(3*t) + np.random.uniform(size=100)

# Create a data frame for the generated data
random_data = pd.DataFrame()
random_data["X"] = t
random_data["Y"] = y1

# Plot the generated data
```




Output:

```
np.linspace() np.sin() np.random.uniform()
```

#### **b) Fit a support vector regression model to the data.**

```
from sklearn.svm import SVR
clf = SVR()
clf = clf.fit(random_data["X"].reshape(-1,1),random_data["Y"])
predictions = pd.DataFrame()
predictions["predY"] = clf.predict(random_data["X"].reshape(-1,1))

plt.scatter(t,y1,color="black")
plt.scatter(t,predictions["predY"],color="r")
plt.show()
```



Output:

SVR


```
predictions["actual"] = random_data["Y"]
predictions["sq_diff"] = (predictions["predY"] - predictions["actual"])**2
print(predictions["sq_diff"].mean())

## 0.20626922373052764
```

### c) Fit a linear regression model to the data.

```
from sklearn import linear_model
linMod = linear_model.LinearRegression()
linMod = linMod.fit(random_data["X"].reshape(-1,1), random_data["Y"])
predictions = pd.DataFrame()
predictions["predY"] = linMod.predict(random_data["X"].reshape(-1,1))

plt.scatter(t,y1,color="black")
plt.scatter(t,predictions["predY"],color="r")
plt.show()
```



Output:

### LinearRegression

```
predictions["actual"] = random_data["Y"]
predictions["sq_diff"] = (predictions["predY"] - predictions["actual"])**2
print(predictions["sq_diff"].mean())

## 4.753313156881628
```

## 6 Model Evaluation & Selection

### 6.1 Evaluate the accuracy of regression models.

#### a) Evaluation on training data.

```
# Notice we are re-using data sets but it is good to re-read the original
# version back into the environment
train = pd.read_csv('/Users/tips_train.csv')
test = pd.read_csv('/Users/tips_test.csv')

# 1. Linear Regression Model
from sklearn.metrics import r2_score
from sklearn import linear_model
linMod = linear_model.LinearRegression()
linMod = linMod.fit(train.drop(["tip"], axis = 1), train["tip"])

# Evaluation on training data
```

```

pred_lin = linMod.predict(train.drop(["tip"], axis = 1))

# Determine coefficient of determination score
r2_lin = r2_score(train["tip"], pred_lin)
print("Linear regression model r^2 score (coefficient of determination): %f"
      % r2_lin)

## Linear regression model r^2 score (coefficient of determination): 0.496730

--

# 2. Random Forest Regression Model
from sklearn.ensemble import RandomForestRegressor
rfMod = RandomForestRegressor(random_state=29)
rfMod = rfMod.fit(train.drop(["tip"], axis = 1), train["tip"])

# Evaluation on training data
pred_rf = rfMod.predict(train.drop(["tip"], axis = 1))

# Determine coefficient of determination score
r2_rf = r2_score(train["tip"], pred_rf)
print("Random forest regression model r^2 score (coefficient of
determination): %f" % r2_rf)

## Random forest regression model r^2 score (coefficient of determination):
0.892204

```

## b) Evaluation on testing data.

```

# 1. Linear Regression Model (linMod)

# Evaluation on testing data
pred_lin = linMod.predict(test.drop(["tip"], axis = 1))

# Determine coefficient of determination score
r2_lin = r2_score(test["tip"], pred_lin)
print("Linear regression model r^2 score (coefficient of determination): %f"
      % r2_lin)

## Linear regression model r^2 score (coefficient of determination): 0.270945

--

# 2. Random Forest Regression Model (rfMod)

# Evaluation on testing data
pred_rf = rfMod.predict(test.drop(["tip"], axis = 1))

# Determine coefficient of determination score
r2_rf = r2_score(test["tip"], pred_rf)

```

```
print("Random forest regression model r^2 score (coefficient of
determination): %f" % r2_rf)

## Random forest regression model r^2 score (coefficient of determination):
0.163330
```

The sklearn metric `r2_score` is only one option for assessing a regression model. Please go [here](#) for more information about other sklearn regression metrics.

## 6.2 Evaluate the accuracy of classification models.

### a) Evaluation on training data.

```
# Notice we are re-using data sets but it is good to re-read the original
# version back into the environment
train = pd.read_csv('/Users/breastcancer_train.csv')
test = pd.read_csv('/Users/breastcancer_test.csv')

# 1. Decision Tree Classification Model
from sklearn import tree
from sklearn.metrics import accuracy_score
treeMod = tree.DecisionTreeClassifier(criterion='entropy', random_state=29)
treeMod = treeMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Evaluation on training data
scored = pd.DataFrame(treeMod.predict(train.drop(["Target"], axis = 1)))
scored["Target"] = train["Target"]

# Determine accuracy score
accuracy_tree = accuracy_score(scored["Target"], scored[0])
print("Decision tree model accuracy: %f" % accuracy_tree)

## Decision tree model accuracy: 1.000000

--

# 2. Random Forest Classification Model
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
rfMod = RandomForestClassifier(random_state=29)
rfMod = rfMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Evaluation on training data
scored = pd.DataFrame(rfMod.predict(train.drop(["Target"], axis = 1)))
scored["Target"] = train["Target"]

# Determine accuracy score
accuracy_rf = accuracy_score(scored["Target"], scored[0])
print("Random forest model accuracy: %f" % accuracy_rf)

## Random forest model accuracy: 0.997487
```

--

```
# 3. Gradient Boosting Classification Model
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score
gbmMod = GradientBoostingClassifier(random_state = 29, learning_rate = .01,
min_samples_leaf = 20, n_estimators = 2500)
gbmMod = gbmMod.fit(train.drop(["Target"], axis = 1), train["Target"])

# Evaluation on training data
scored = pd.DataFrame(gbmMod.predict(train.drop(["Target"], axis = 1)))
scored["Target"] = train["Target"]

# Determine accuracy score
accuracy_gbm = accuracy_score(scored["Target"], scored[0])
print("Gradient boosting model accuracy: %f" % accuracy_gbm)

## Gradient boosting model accuracy: 1.000000
```

## b) Evaluation on testing data.

```
# 1. Decision Tree Classification Model (treeMod)

# Evaluation on testing data
scored = pd.DataFrame(treeMod.predict(test.drop(["Target"], axis = 1)))
scored["Target"] = test["Target"]

# Determine accuracy score
accuracy_tree = accuracy_score(scored["Target"], scored[0])
print("Decision tree model accuracy: %f" % accuracy_tree)

## Decision tree model accuracy: 0.947368
```

--

```
# 2. Random Forest Classification Model (rfMod)

# Evaluation on testing data
scored = pd.DataFrame(rfMod.predict(test.drop(["Target"], axis = 1)))
scored["Target"] = test["Target"]

# Determine accuracy score
accuracy_rf = accuracy_score(scored["Target"], scored[0])
print("Random forest model accuracy: %f" % accuracy_rf)

## Random forest model accuracy: 0.964912
```

--

```
# 3. Gradient Boosting Classification Model (gbmMod)

# Evaluation on testing data
```



```

scored = pd.DataFrame(gbmMod.predict(test.drop(["Target"], axis = 1)))
scored["Target"] = test["Target"]

# Determine accuracy score
accuracy_gbm = accuracy_score(scored["Target"], scored[0])
print("Gradient boosting model accuracy: %f" % accuracy_gbm)

## Gradient boosting model accuracy: 0.976608

```

Note: The sklearn metric `accuracy_score` is only one option for assessing a classification model. Please go [here](#) for more information about other sklearn classification metrics.

## 6.3 Evaluation with cross validation.

### a) KFold

```

# Notice we are using a new data set that need to be read into the
# environment
breastcancer = pd.read_csv('/Users/breastcancer.csv')

from sklearn import model_selection
from sklearn.ensemble import RandomForestClassifier

X = breastcancer.drop(["Target"], axis = 1)
Y = breastcancer["Target"]

kfold = model_selection.KFold(n_splits = 5, random_state = 29)
model = RandomForestClassifier(random_state = 29)
results = model_selection.cross_val_score(model, X, Y, cv = kfold)

print("Accuracy: %.2f%% +/- %.2f%%" % (results.mean()*100,
                                       results.std()*100))

## Accuracy: 94.38% +/- 2.39%

```

### b) ShuffleSplit

```

shuffle = model_selection.ShuffleSplit(n_splits = 5, random_state = 29)
model = RandomForestClassifier(random_state = 29)
results = model_selection.cross_val_score(model, X, Y, cv = shuffle)

print("Accuracy: %.2f%% +/- %.2f%%" % (results.mean()*100,
                                       results.std()*100))

## Accuracy: 95.09% +/- 0.70%

```

## Appendix

### 1 Built-in Python Data Types

- `Boolean`

#### Numeric types:

- `int`
- `long`
- `float`
- `complex`

#### Sequences:

- `str`
- `bytes`
- `byte array`
- `list`
- `tuple`

#### Sets:

- `set`
- `frozen set`

#### Mapping:

- `dictionary`

### 2 Python Plotting Packages

**Bokeh**

**PyPlot**

**Seaborn**

---

## Alphabetical Index

### Array

A NumPy array is a data type implemented by the [NumPy](#) package in which the elements of the array are all of the same type. Please see the following example of array creation and access:

```
import numpy as np
my_array = np.array([1, 2, 3, 4])
print(my_array)

## [1 2 3 4]

print(my_array[3])

## 4
```

For more information, please see [NumPy Arrays](#).

---

## Bokeh

[Bokeh](#) is a Python package which is useful for interactive visualizations and is optimized for web browser presentations.

---

## Boolean

A [Boolean](#) value is either True or False, and represents the truth of an expression or statement.

---

## Bytes & Byte arrays

A [byte](#) is a sequence of integers which is immutable, whereas a [byte array](#) is its mutable counterpart.

---

## complex

A [complex number](#) includes a real part and an imaginary part, both of which are floating point numbers.

---

## Data Frame

A [Pandas Data Frame](#) is a two-dimensional tabular structure with labeled axes (rows and columns), where data observations are represented by rows and data variables are represented by columns.

---

## datetime

The [datetime](#) Python module includes tools for manipulating data and time objects.

---

## Decimal

[Decimal](#) is a Python package which provides tools for decimal floating point arithmetic.

---

## Dictionary

A [dictionary](#) is an associative array which is indexed by keys which map to values. Therefore, a dictionary is an unordered set of key:value pairs where each key is unique. Please see the following example of dictionary creation and access:

```
import pandas as pd
student = pd.read_csv('/Users/class.csv')
for_dict = pd.concat([student["Name"], student["Age"]], axis = 1)
class_dict = for_dict.set_index('Name').T.to_dict('list')
print(class_dict.get('James'))

## [12]
```

---

## float

A [float](#) is a decimal point number.

---

## int

An [int](#) is a natural number. In Python, you can convert to an int from a float by using the `int()` function. Python stores ints with at least 32 bits of precision.

---

## List

A [list](#) is a sequence of comma-separated objects that need not be of the same type. Please see the following example of list creation and access:

```
list1 = ['item1', 102]
print(list1)

## ['item1', 102]
```

---

```
print(list1[1])
```

```
## 102
```

Python also has what are known as "**Tuples**", which are immutable lists created in the same way as lists, except with paranthesis instead of brackets.

---

## Long

A **long** is a type of integer with unlimited precision. In Python, you can convert to a long using the `long()` function.

---

## NumPy

**NumPy** is a Python package which is useful for scientific and mathematical computing.

---

## pandas

**pandas** is a Python package which is useful for working with data structures and performing data analysis.

---

## PyPlot

**PyPlot** is a Python package which is useful data plotting and visualization.

---

## Seaborn

**Seaborn** is another Python package which is useful for data plotting and visualization. In particular, Seaborn includes tools for drawing attractive statistical graphics.

---

## Series

A **Pandas Series** is a one-dimensional data frame, which is also called an array in R. Please see the following example of Series creation and access:

```
import pandas as pd
my_array = pd.Series([1, 3, 5, 9])
print(my_array)
```

```
## 0      1
## 1      3
## 2      5
## 3      9
## dtype: int64

print(my_array[1])

## 3
```

---

## Sets & Frozen Sets

A set is a unordered collection of immutable objects. The difference between a [set](#) and a [frozen set](#) is that the former is mutable, while the latter is immutable. Please see the following example of set and frozen set creation and access:

```
s = set(["1", "2", "3"])
print(s)
# s is a set, which means you can add or delete elements from s

## {'3', '1', '2'}

fs = frozenset(["1", "2", "3"])
print(fs)
# fs is a frozenset, which means you cannot add or delete elements from fs

## frozenset({'2', '3', '1'})
```

---

## sklearn

scikit-learn, or more commonly known as [sklearn](#), is a Python package which is useful for basic and advanced data mining, machine learning, and data analysis. sklearn includes tools for classification, regression, clustering, dimensionality reduction, model selection, and data pre-processing.

---

## str

[Strings](#) are a list of characters, though characters are not a type in Python, but rather a string of length 1. Strings are indexable like arrays. Please see the following example of String creation and access:

```
s = 'My first string!'
print(s)

## My first string!
```

```
print(s[5])  
## r
```

Please go [here](#) for more information on the `str()` function.

---

## **sub**

`sub` is a function of the `re` Python package useful for replacing a pattern in a string.

---

For more information on Python packages and functions, along with helpful examples, please see [Python](#).